

Performance and Robustness Benchmarking Across Battery State-of-Charge Estimator Classes for Embedded Applications

Anonymous

Abstract

Embedded battery-state estimators require validation under disturbed inputs and the inference semantics used on resource-constrained hardware. This study presents a causal, cell-disjoint benchmark of four State-of-Charge (SOC) approaches: fixed-capacity Coulomb counting as a direct-measurement model (DM), Coulomb counting with data-driven State-of-Health adaptation as a hybrid direct-measurement model (HDM), a second-order resistor-capacitor observer as a hybrid equivalent-circuit model (HECM), and a gated recurrent unit structurally pruned from 96 to 67 hidden units as the data-driven model (DD). Public lithium iron phosphate ageing data cover multiple operating and health conditions. Robustness evaluation addresses nominal accuracy, corrupted or unavailable measurements, and recovery from initialization errors. Hardware evaluation on an STM32H753 measures numerical equivalence, latency, flash, and peak runtime memory. For DD, exact rolling-window recomputation, continuous state forwarding, and periodic state reset are compared. Nominal cell-macro mean absolute errors are 0.0690, 0.0412, 0.0337, and 0.0258 for DM, HDM, HECM, and DD. DD yields the lowest nominal error and leads several disturbance scenarios, whereas HECM shows the strongest recovery. All implementations reproduce their software references, fit available memory, and meet the one-second sampling deadline. Compared with rolling-window recomputation, continuous-state DD reduces median latency from approximately 724 ms to 0.425 ms and peak memory from 67.3 KiB to 4.1 KiB while retaining similar nominal error. Periodic resets retain these resource benefits but cause pronounced transients. The

results demonstrate that estimator choice, model compression, and recurrent-state handling must be evaluated jointly with accuracy, robustness, recovery, and embedded cost.

Keywords: State-of-charge estimation, Battery management system, Recurrent neural network, Robustness benchmarking, Embedded artificial intelligence, Model compression

1 Nomenclature

2 *Acronyms*

3	BMS	Battery management system
4	SOC	State of charge
5	SOH	State of health
6	DM	Direct-measurement model
7	HDM	Hybrid direct-measurement model
8	HECM	Hybrid equivalent-circuit model
9	DD	Data-driven neural model
10	ECM	Equivalent circuit model
11	EKF	Extended Kalman filter
12	MAE	Mean absolute error
13	RMSE	Root-mean-square error
14	ADC	Analog-to-digital converter
15	CV	Constant-voltage phase
16	OCV	Open-circuit voltage
17	GRU	Gated recurrent unit
18	LSTM	Long short-term memory
19	LFP	Lithium iron phosphate

20 *Symbols*

21	t	Time
22	$I(t)$	Current signal at time t
23	$U(t)$	Voltage signal at time t
24	$T(t)$	Temperature signal at time t
25	$\widehat{\text{SOC}}(t)$	Estimated state of charge at time t
26	$\widehat{\text{SOH}}(t)$	Estimated state of health at time t
27	$\text{SOC}_{\text{ref}}(t)$	Dataset SOC ground truth at time t
28	$\text{SOH}_{\text{ref}}(t)$	Reference SOH target at time t
29	g_I	Multiplicative current-gain error
30	ΔI	Additive current offset
31	ΔU	Additive voltage offset
32	ΔT	Additive temperature offset
33	σ_I	Standard deviation of injected current noise
34	σ_U	Standard deviation of injected voltage noise
35	σ_T	Standard deviation of injected temperature noise
36	Q_c	Online cumulative charge state used by the SOC models
37	C_{nom}	Nominal capacity
38	C_{eff}	Effective capacity used inside the estimator
39	U_{max}	Upper voltage threshold used for full-charge detection
40	U_{tol}	Voltage tolerance used in the full-charge condition
41	U_1, U_2	Polarization voltages of the first and second RC branch
42	ΔMAE	Error increase relative to baseline

43 **1. Introduction**

44 *1.1. Motivation and practical relevance*

45 Reliable battery state estimation is a core BMS function because SOC and
46 SOH determine safety margins, usable energy, operating limits, charge con-
47 trol, and degradation-aware system management in both mobile and stationary

lithium-ion battery applications [1, 2, 3]. Practical deployment quality is not defined by clean-condition accuracy alone, because real BMS operation is exposed to sensor gain and offset errors, additive noise, initialization uncertainty after restart, missing samples, irregular sampling, communication interruptions, and transient signal spikes caused by measurement-chain limitations or disturbances in the surrounding system [4, 5, 6]. The relevant engineering question is therefore not only whether an estimator achieves low MAE under nominal measurements, but also whether it converges back to a reliable state after disturbed initialization and whether it remains stable under physically plausible measurement corruption.

As summarized conceptually in Figure 1, the present benchmark treats deployment-facing estimator quality as a three-dimensional problem consisting of nominal accuracy, recovery after state inconsistency, and robustness against disturbed measurements and signal-integrity faults, embedded within broader constraints on memory and execution time. The central question is how four structurally different SOC-estimation classes, represented by the implementations studied here, respond to common failure mechanisms and microcontroller constraints, and whether their nominal ranking remains informative under those conditions. This distinction is essential because one implementation can be accurate under clean measurements yet difficult to re-anchor, sensitive to a specific input fault, or too costly under its trained inference semantics.

1.2. State of the art in battery state estimation

Battery-state estimators are commonly grouped into four broad families [7, 1, 2, 3]. These are direct or integration-based methods, model-based methods using equivalent-circuit or electrochemical models, hybrid methods combining physical structure and learned components, and data-driven neural models. Direct-measurement approaches, typically based on Coulomb counting, are attractive because they are simple, computationally inexpensive, and easy to deploy. However, they remain structurally sensitive to current-gain error, missing data, and initialization errors because they lack an internal correction mechanism [7, 4].




Figure 1: Conceptual overview of deployment-oriented requirements for embedded BMS state estimation. The broader requirement space includes implementation effort, hardware constraints, and estimator performance, whereas the present study focuses on the performance branch and evaluates baseline accuracy, convergence behavior, and robustness under measurement disturbances.

78 Model-based approaches, especially ECM- and observer-based estimators, of-
 79 fer physical interpretability and voltage-based correction capability, but their
 80 practical behavior depends strongly on the quality of model parametrization,
 81 operating-point observability, and measurement quality [8, 2, 4].

82 Hybrid estimators seek to combine physical state propagation with learned
 83 adaptation of capacity, parameters, or auxiliary states, which has made joint
 84 SOC/SOH estimation an active research direction [9, 10, 11]. This combination
 85 can improve baseline calibration, while it does not necessarily eliminate struc-
 86 tural sensitivity to measurement disturbances. Data-driven approaches based on
 87 recurrent or sequence models can exploit nonlinear cross-relations between cur-
 88 rent, voltage, temperature, and derived online features. Recent studies have also
 89 demonstrated their embedded feasibility through TinyML-oriented implemen-
 90 tations, pruning, and quantization [12, 13, 14, 15, 3]. For recurrent estimators,

91 however, deployment cost is determined not only by parameter count but also
92 by how input history and hidden state are represented during inference.

93 *1.3. Performance and embedded deployment gap*

94 A large share of the SOC estimation literature still emphasizes clean-data ac-
95 curacy, average error metrics, or computational efficiency, while comparatively
96 fewer studies analyze how estimators behave under realistic measurement faults
97 and timing corruption [1, 13, 14, 15]. Existing robustness studies often remain
98 confined to a single model family, for example observer-based robustness against
99 modeling and measurement errors [4] or robust estimation concepts for specific
100 BMS architectures under noise and uncertainty [5]. Cross-family comparisons
101 under the same disturbance protocol and comparable embedded constraints re-
102 main limited. Benchmarking studies in the broader BMS context underline the
103 importance of reproducible scenario design and deployment-relevant validation,
104 but they do not close the gap between clean-data estimator comparison and
105 measurement-level evaluation across fundamentally different estimator classes
106 [6]. Embedded studies also tend to report one exported network without iso-
107 lating how recurrent execution semantics affect accuracy, latency, and memory.
108 This gap is especially relevant in BMS design, where a compact architecture can
109 still be inefficient if complete input windows are repeatedly evaluated, while an
110 efficient stateful implementation can fail if its hidden state is reset inappropri-
111 ately.

112 *1.4. Contribution and study objectives*

113 This work addresses the gap by examining whether clean-data accuracy pre-
114 dicts behavior under measurement and signal-integrity disturbances, which up-
115 date mechanisms govern error growth and recovery, and how these properties
116 interact with embedded execution cost. Four deployment-oriented representa-
117 tives are evaluated under one causal protocol on cell-disjoint LFP aging data
118 covering multiple load and SOH states. The campaign combines physically in-
119 terpretable measurement, timing, availability, quantization, and initialization

disturbances with a common recovery criterion, cell-level uncertainty analysis, a paired SOH-input ablation, and STM32 measurements of the isolated SOC cores.

The artificial-intelligence contribution is an end-to-end evaluation of a compact recurrent SOC estimator within this common engineering benchmark. A structurally pruned and fine-tuned GRU is compared with conventional estimator classes and is then executed using exact rolling-window recomputation, continuous hidden-state forwarding, and periodic state reset. Keeping the trained weights, preprocessing, and numerical implementation fixed across these modes isolates the effect of recurrent-state handling on accuracy, latency, and memory. The resulting evidence connects learned temporal representation and embedded implementation semantics with the fault profile and resource constraints of the target BMS application.

2. Methodology

2.1. Estimator representatives and class boundaries

Figure 2 summarizes the causal benchmark chain, from measured signals and online feature reconstruction to disturbance injection and global and local evaluation. The four estimators were selected because their dominant SOC update and correction mechanisms differ. Table 1 relates these mechanisms to the broader design space and specifies the implementation evaluated for each class. Learned corrections, higher-order dynamics, richer sensing, and more complex sequence architectures extend these class envelopes while also changing deployment cost.

Table 1: Estimator-class envelope and implementation scope. Literature in the second column supports the defining mechanisms and broader capability ranges, while the last column specifies the representative evaluated in this study.

Class	Defining mechanism and broader capability range	Representative used in this study
Direct measurement	Causal integration of measured current using an assumed capacity. Extensions can include coulombic efficiency, capacity adaptation, OCV or end-point anchoring, plausibility logic, and reset thresholds [7, 13].	Fixed-capacity Coulomb counting with a sustained full-charge CV reset and no voltage-residual correction
Hybrid direct measurement	Integration core augmented by an estimated or learned state or parameter. Extensions can include learned capacity or efficiency, residual correction, multi-sensor fusion, and adaptive anchoring [10, 11].	Same Coulomb-counting and reset logic as DM, with effective capacity supplied by the shared causal LSTM-SOH estimate
ECM observer	Coulomb-counted SOC propagation combined with a voltage model and feedback correction. Extensions can include higher RC order, hysteresis, thermal states, operating-point maps, and different nonlinear observers [7, 4, 2].	Second-order RC model with EKF correction, charge/discharge SOC-SOH lookup surfaces, and shared causal SOH input

Table 1: Estimator-class envelope and implementation scope (continued).

Class	Defining mechanism and broader capability range	Representative used in this study
Data driven	Learned mapping from measured signals or their history to SOC. Implementations can use feed-forward, convolutional, recurrent, attention-based, physics-informed, raw-signal, or engineered-feature models [12, 15, 13].	Structurally pruned rolling-window GRU using voltage, current, temperature, SOH, Q_c , derivatives, and Δt

143 All SOH-dependent representatives, namely HDM, HECM, and DD, receive
 144 the same causal LSTM-SOH trace so that differences arise from how the SOC
 145 branch consumes that context. DM is the low-complexity integration reference.
 146 HDM isolates the effect of learned capacity adaptation without adding volt-
 147 age feedback. HECM adds model-based voltage-residual correction to current-
 148 driven state propagation. DD represents compact temporal learning. The GRU
 149 was selected because SOC and disturbance recovery depend on recent signal his-
 150 tory while its gated single-state structure is lighter than a comparable LSTM.
 151 Because DD also consumes the engineered Q_c feature, it is a data-driven se-
 152 quence estimator with physics-informed inputs rather than a raw-signal-only
 153 neural model.

154 2.2. SOC estimation methods

155 *Direct-measurement model (DM)*

156 The direct-measurement SOC formulation used in DM, HDM, and as an
 157 explicit auxiliary feature in DD follows the same online cumulative-charge state,
 158 denoted here uniformly as Q_c . It is updated causally from the measured current
 159 stream as

$$Q_c(k) = Q_c(k-1) + \frac{I_k \Delta t_k}{3600 \text{ s h}^{-1}}, \quad (1)$$




Figure 2: Methodology overview for the causal benchmark chain. Measured battery signals are converted into online auxiliary features, processed by representative estimator implementations, exposed to measurement and initialization interventions, and finally evaluated through global and local performance metrics.

where I_k is given in ampere and Δt_k in seconds, so that the factor 3600 s h^{-1} converts the increment to ampere-hours. The sign convention is fixed by the benchmark implementation and the state is reset to $Q_c = 0$ after a sustained constant-voltage full-charge event. In the DM, this same charge state is mapped directly to SOC through the fixed nominal capacity

$$\widehat{\text{SOC}}_{\text{DM}}(k) = \text{clip}\left(1 + \frac{Q_c(k)}{C_{\text{nom}}}, 0, 1\right). \quad (2)$$

Hybrid direct-measurement model (HDM)

The HDM preserves exactly the same Coulomb-counting core and differs only in the capacity term. The SOH estimate rescales the usable capacity as

$$C_{\text{eff}}(k) = C_{\text{nom}} \cdot \widehat{\text{SOH}}(k), \quad (3)$$

which yields

$$\widehat{\text{SOC}}_{\text{HDM}}(k) = \text{clip}\left(1 + \frac{Q_c(k)}{C_{\text{eff}}(k)}, 0, 1\right). \quad (4)$$

This formulation makes the relation between DM and HDM explicit: both use the same online charge state Q_c , while only the denominator changes through online SOH support. In both direct-measurement variants, a sustained near-maximum voltage condition is interpreted as a constant-voltage full-charge event. If the voltage remains above $U_{\text{max}} - U_{\text{tol}}$, where U_{tol} denotes the admissible voltage margin below the upper threshold U_{max} , for the configured CV duration,

the internal charge state is reset so that $Q_c = 0$ and therefore $\widehat{\text{SOC}} = 1$. This reflects the practical assumption that a completed CV phase provides a physically anchored full-charge reference.

Hybrid equivalent-circuit model (HECM)

The hybrid ECM model is implemented as a second-order RC model with the state vector

$$\mathbf{x}_k = \begin{bmatrix} \widehat{\text{SOC}}_k & U_{1,k} & U_{2,k} \end{bmatrix}^\top, \quad (5)$$

where U_1 and U_2 denote the polarization voltages of the two RC branches. The prediction step propagates the state according to

$$\mathbf{x}_k^- = \mathbf{A}_k \mathbf{x}_{k-1} + \mathbf{b}_k I_{k-1}, \quad (6)$$

with

$$\mathbf{A}_k = \text{diag}\left(1, e^{-\Delta t_k/\tau_1}, e^{-\Delta t_k/\tau_2}\right), \quad \mathbf{b}_k = \begin{bmatrix} \eta \Delta t_k / C_b \\ R_1(1 - e^{-\Delta t_k/\tau_1}) \\ R_2(1 - e^{-\Delta t_k/\tau_2}) \end{bmatrix}, \quad (7)$$

where $C_b = C_0 \widehat{\text{SOH}}$ is the SOH-scaled available capacity. The corrected terminal-voltage estimate is written as

$$\widehat{U}_k = \text{OCV}(\widehat{\text{SOC}}_k, \widehat{\text{SOH}}_k, I_k) + U_{1,k} + U_{2,k} + R_i I_k, \quad (8)$$

and the Kalman update uses the residual between $\widehat{U}_k$ and the measured terminal voltage to correct the internal state. In the present implementation, the state propagation is driven by the previous current sample I_{k-1} , whereas the measurement equation uses the current sample I_k available at the correction step. The ECM parameters are not constant. The implementation interpolates R_i , R_1 , R_2 , τ_1 , τ_2 , OCV, and dOCV/dSOC from a preidentified lookup table over SOC and SOH, with separate charge and discharge parameter surfaces. The HECM is implemented as a two-RC observer with SOH-dependent parameter scheduling.

195 *Data-driven SOC model (DD)*

196 The data-driven SOC model is a GRU-based recurrent estimator with the
 197 feature vector $\{U \text{ [V]}, I \text{ [A]}, T \text{ [}^\circ\text{C]}, \text{SOH}, Q_c, dU/dt, dI/dt, \Delta t\}$. Its single-
 198 layer recurrent core is followed by an MLP head with one hidden layer of size
 199 96 and a sigmoid output. During optimization, the recurrent core is exposed to
 200 causal sequence segments of length $L_{\text{SOC}} = 2024$ samples. Within the feature
 201 set, Q_c is the online cumulative-charge state defined in (1). It supplies the
 202 neural estimator with the same charge-throughput information that drives DM
 203 and HDM. The derivative channels are calculated online as $dI/dt(k) = (I_k -$
 204 $I_{k-1})/\Delta t_k$ and $dU/dt(k) = (U_k - U_{k-1})/\Delta t_k$. They represent local transients,
 205 while Δt_k allows the model to distinguish signal changes from changes in the
 206 sampling interval.

207 The deployed DD architecture is obtained by structurally pruning a GRU
 208 with 96 recurrent units. Removing 29 complete units reduces the hidden-state
 209 size to 67, corresponding to a 30.2% reduction in recurrent width, and removes
 210 the associated recurrent and downstream connections. The resulting matri-
 211 ces remain dense and can therefore be evaluated directly by the hand-written
 212 floating-point C kernels without sparse-matrix support. Fine-tuning follows
 213 pruning, and all training, pruning, and model-selection decisions use only the
 214 declared training and validation data. The holdout cells are reserved for the sub-
 215 sequent robustness and hardware evaluation. Because the cross-class campaign
 216 evaluates this fixed compact architecture rather than an additional unpruned
 217 DD control, pruning defines the deployment-ready neural representative and is
 218 not treated as a separate benchmark factor.

219 *2.3. SOH estimation method*

220 The shared SOH estimator is an LSTM-based sequence-to-sequence model
 221 that operates on hourly aggregated online features derived from $\{U \text{ [V]}, I \text{ [A]},$
 222 $T \text{ [}^\circ\text{C]}, \text{EFC}, Q_c\}$, where the aggregation uses the statistics mean, standard
 223 deviation, minimum, and maximum for each feature [16, 17, 18, 19]. This
 224 transforms the five base channels into a 20-dimensional hourly feature vector.

Architecturally, the SOH network first projects the aggregated input through a two-layer feature embedding block with size 128, layer normalization, GELU activations, and dropout, then processes the sequence with a two-layer unidirectional LSTM of hidden size 112, and finally maps the recurrent output through three residual MLP blocks and a prediction head with hidden size 160 to obtain the SOH estimate for each hourly step.

During optimization, the SOH estimator is exposed to causal hourly sequence segments, and benchmark execution processes one hourly aggregate at a time with causal hidden- and cell-state forwarding. This allows the estimator to accumulate degradation context without access to future information. The first hourly output is anchored with the configured initial SOH value before subsequent hourly predictions are taken directly from the recurrent model output. The SOH prediction is updated at hourly cadence and held piecewise constant between update points. This online update structure matches the gradually evolving capacity fade in the dataset and supplies the latest SOH context to the faster SOC estimators. In the HDM, the predicted SOH enters the SOC computation through the effective capacity relation $C_{\text{eff}} = C_{\text{nom}} \cdot \widehat{\text{SOH}}$, so the method remains structurally close to direct measurement while replacing the fixed-capacity assumption. In the HECM, the predicted SOH affects the capacity scaling and the lookup-based ECM parametrization, while in the DD model it is provided as an explicit SOC input feature that contextualizes the current measurement trajectory.

A paired diagnostic comparison is used to distinguish errors produced within each SOC branch from errors inherited through the estimated SOH input. For selected scenarios, the causal LSTM-SOH trace is replaced by the offline dataset ground truth while the SOC implementation, disturbed measurements, evaluation window, and error calculation remain unchanged. This substitution is not a deployable estimator configuration. It isolates how strongly each SOC branch depends on the accuracy and characteristics of the supplied health information.

254 2.4. Model and target-hardware preparation

255 The software benchmark uses trained floating-point model parameters and
256 input scalers fixed before evaluation on the independent holdout test set. Neural
257 training, hyperparameter selection, scaler fitting, pruning, and model selection
258 use the declared training and validation sets. The HPPC-derived HECM lookup
259 surfaces were identified from the same model-development sets and fixed without
260 retuning before benchmark evaluation. Figure A.25 and Table A.4 document
261 the complete cell-disjoint split. The compact SOC-GRU has hidden size 67
262 and the shared SOH-LSTM has hidden size 112. The continuous high-load test
263 trajectory in Figure 6 is analyzed separately as a mechanism-focused lifecycle
264 diagnostic.

265 Three execution modes are implemented for the fixed pruned DD network.
266 Exact rolling-window execution resets the hidden state and recomputes the lat-
267 est 2024 samples for every output, reproducing the sequence-to-one semantics
268 used during training and in the primary robustness benchmark. Continuous-
269 state execution processes each new sample once and forwards the resulting hid-
270 den state to the next update. Periodic-reset execution uses the same incremen-
271 tal calculation but clears the hidden state after every 2024 updates. All three
272 modes use identical weights, feature scaling, and floating-point kernels. Their
273 comparison therefore isolates the resource and estimation consequences of se-
274 quence recomputation and recurrent-state retention rather than differences in
275 network parameters.

276 The target-hardware experiment evaluates the isolated SOC cores on an
277 STM32H753 Cortex-M7 running at 480 MHz. The processor is mounted on a
278 NUCLEO-H753ZI development board. To separate SOC implementation cost
279 from the shared SOH service, the board receives the precomputed causal SOH
280 trace used by the software reference. The SOH-LSTM itself is not executed on
281 the microcontroller. Numerical equivalence and inference latency are measured
282 directly on the board. Execution time is recorded for every valid inference with
283 the processor-integrated hardware cycle counter, and each cell sequence is re-
284 played three times. Flash occupancy is calculated from the compiled firmware

285 regions that must be stored in the microcontroller’s non-volatile program mem-
286 ory.

287 Peak runtime RAM combines statically allocated variable storage, maximum
288 dynamic-memory allocation, and maximum temporary call-stack demand. The
289 static contribution is obtained from the compiled firmware, while dynamic allo-
290 cation is monitored during execution. Stack demand is measured with a high-
291 water-mark pattern during representative on-device replays that exercise the
292 complete estimator path. For DD, this includes valid inference after its rolling
293 input window has been filled. An unused compilation-time memory reserve is
294 excluded from the reported demand. Predictions are compared with both the
295 software implementation and dataset SOC. Figure 19 summarizes the resulting
296 memory requirements of the isolated SOC implementations with SOH supplied
297 as an external causal trace.

298 *2.5. Robustness benchmark design*

299 The primary cross-model benchmark manipulates current, voltage, tempera-
300 ture, sample availability, and sampling time, all of which are measured or recon-
301 structed online by a BMS. A separate initialization branch acts on deployment-
302 accessible estimator states. Parameter mismatch is excluded from this common
303 ranking and examined through the supplementary HECM analysis defined in
304 Section 2.7. Figure 3 groups the primary scenarios into input disturbances,
305 initialization errors, and signal-integrity faults, while Table 2 defines their im-
306 plementation [6].

307 The disturbance magnitudes follow three practical levels of justification:
308 deployment-level sensing limits, adverse measurement corruption, and signal-
309 integrity stress conditions [20, 21, 22, 23, 24, 25]. Measurement disturbances
310 are separated by mechanism. Random sensor noise represents zero-mean fluc-
311 tuations in current, voltage, and temperature. Systematic sensor errors com-
312 prise multiplicative current-gain error and additive offsets in current, voltage,
313 and temperature. Voltage spikes and ADC quantization complete the input-
314 disturbance branch. The remaining scenarios address estimator initialization,




Figure 3: Taxonomy of the disturbance scenarios used in the performance benchmark, grouped into input disturbances, initialization errors, and signal-integrity faults.

sample availability, and sampling-time integrity. Together they form a controlled and reproducible benchmark matrix for exposing estimator-specific failure mechanisms under a common protocol.

Table 2: Disturbance configurations used in the benchmark and their practical rationale. Exact numeric values are fixed benchmark settings. The cited sources motivate disturbance type and magnitude range.

Scenario	Benchmark setup	Practical origin / rationale
ADC quantization	$\Delta I = 0.01$ A, $\Delta U = 0.005$ V, and $\Delta T = 0.5$ °C	ADC steps, sensor scaling, rounding, and coarse deployment-level discretization [20, 21, 26]
Current gain error	$I_{\text{meas}} = I(1 + g_I)$ with $g_I = \pm 0.5\%$, $\pm 1.5\%$, and $\pm 3.0\%$. Both signs are evaluated at each magnitude	Sensor sensitivity error, scale-factor miscalibration, drift, and upper bound as stress case [21, 22, 23]
Current noise	Gaussian with $\sigma_I = 0.02$ A and 0.10 A	Sensing-chain noise, EMI pickup, and low and high sensor-noise levels [4, 5]
Voltage noise	Gaussian with $\sigma_U = 0.01$ V	Wiring noise, front-end noise, poor filtering, and EMI [4, 20, 24]

Table 2: Disturbance configurations used in the benchmark and practical rationale (continued).

Scenario	Benchmark setup	Practical origin / rationale
Temperature noise	Gaussian with $\sigma_T = 1.0^\circ\text{C}$	Sensor tolerance, conversion noise, thermal gradients, and sensor-to-cell mismatch [26, 11]
Current offset	Additive $I_{\text{meas}} = I + \Delta I$ with $\Delta I = \pm 0.05\text{ A}$. Both signs are evaluated	Zero-current calibration residual and systematic current-sensor offset [21, 22, 23]
Voltage offset	Additive $U_{\text{meas}} = U + 0.02\text{ V}$	Front-end offset, calibration residual, wiring contribution, and adverse sensing stress
Temperature offset	Additive $T_{\text{meas}} = T + 3^\circ\text{C}$	Sensor placement error, calibration residual, and cell-to-sensor thermal mismatch
Initial SOC error	Initial mismatch of -0.10 SOC, with DD perturbed through online Q_c	Restart, storage, missing recent history, and weak LFP OCV anchoring [4, 8]
Missing samples	Every 50th sample frozen and separate random 2% loss	Packet loss, blocked updates, logger dropouts, and periodic and random data freeze [5, 25]
Irregular sampling	Uniform jitter: $\pm 0.1\text{ s}$, $\pm 0.5\text{ s}$, and $\pm 0.9\text{ s}$ around nominal 1 s grid	Scheduling jitter, bus latency, and asynchronous acquisition [6, 5, 25]
Burst dropout	One frozen 3600 s interval selected by maximum absolute net charge, with at least 12 h before and 24 h after the interval	Communication interruption, frozen task, stuck data stream, and measurement-only severe case [6, 5]
Voltage spikes	Single-sample spikes of $\pm 0.20\text{ V}$ every 1000 samples (1000 s)	EMI, switching disturbance, and corrupted voltage samples [24]

318 2.6. Metrics and evaluation criteria

319 For deployment-oriented interpretation, the benchmark results are synthe-
320 sized along three performance dimensions: nominal accuracy under baseline con-
321 ditions, recovery after initialization mismatch, and robustness under sustained
322 or transient measurement corruption and signal unavailability. Disturbance-
323 related global performance is quantified through MAE, RMSE, bias, maximum
324 error, and the 95th percentile absolute error, because these metrics provide a
325 compact view of average performance, systematic deviation, and tail behavior
326 across the selected evaluation windows. Local disturbance performance is an-
327 alyzed through scenario-specific metrics such as jump counts, output variance,
328 absolute-error variance, and excess rolling MAE, because several disturbances do
329 not primarily manifest as window-wide MAE collapse but as short-term output
330 fluctuations, delayed drift, or slow re-synchronization. Whenever a disturbance
331 mask exists, the analysis further separates pre-disturbance, disturbed, and post-
332 disturbance behavior in order to quantify disturbance-window error growth and
333 residual post-disturbance error.

334 ADC quantization, current-gain error, additive sensor offsets, voltage and
335 temperature noise, missing samples, and irregular sampling are interpreted pri-
336 marily through global metrics. Burst dropout, voltage spikes, current-noise
337 detail, and initialization mismatch additionally use dedicated local evidence
338 because their most relevant effects are episodic. The canonical initialization-
339 recovery analysis compares each perturbed estimator trajectory with its cell-,
340 window-, model-, SOH-, and seed-matched correctly initialized trajectory. The
341 resulting paired trajectory difference is the recovery signal, while absolute error
342 against dataset SOC remains the accuracy measure. A first qualifying entry
343 occurs when the paired difference enters a 0.02 SOC band for at least 300 s.
344 Persistent recovery is the first subsequent point after which the difference re-
345 mains within this band for the rest of the 24-h observation horizon. Runs
346 without this endpoint are right-censored. First-entry time, later departure from
347 the band, censored and relapse fractions, and excess-error area under the curve
348 provide complementary recovery diagnostics. The protocol maps a -0.10 SOC-

equivalent mismatch into the coulomb-counting SOC initialization for DM and HDM, the EKF SOC state for HECM, and the capacity-equivalent Q_c offset for DD. Because DD first requires a causal input context, recovery is evaluated for every estimator from the first time at which all four provide a valid output. Recovery times remain referenced to the original intervention. Endpoints already satisfied at the common evaluation start are treated as left-censored and recorded as conservative upper bounds. The comparison therefore measures observable re-anchoring after deployment-accessible, estimator-specific initialization interventions.

Each holdout cell is treated as one independent statistical unit. Protocol-defined SOH windows and random seeds are first summarized within each cell, after which the six cells receive equal macro weight. Hierarchical bootstrap resampling with seeds nested inside cells and 10,000 repetitions provides 95% confidence intervals. Scenario effects are evaluated as paired, cell-matched differences from baseline, and estimator pairs are compared using exact two-sided sign-flip tests, raw mean and median differences, paired standardized effect sizes, and Holm-adjusted p -values. Holm correction is applied within each reported scenario family. With six independent cells, effect magnitudes and confidence intervals provide the primary evidence and p -values support their interpretation. The full machine-readable tests accompany the result tables, while the baseline pair tests are reproduced in the Appendix.

To ensure like-for-like comparisons, all global accuracy, robustness, stratified, confidence-interval, and paired-test calculations begin when the rolling DD first provides a valid estimate. The resulting common evaluation mask gives every estimator the same source interval and the same number of scored observations.

The decision synthesis provides a relative summary of the four implementations. Accuracy combines baseline MAE, RMSE, and 95th-percentile error. Robustness assigns equal weight to eight evaluated disturbance families: random sensor noise, current-gain error, additive sensor offsets, quantization, missing samples, timing jitter, burst dropout, and voltage spikes. The current-gain and current-offset components use the adverse direction from their matched signed

sweeps. Non-positive ΔMAE values contribute zero disturbance penalty. Levels within a family are averaged before the family scores are combined, giving every disturbance family equal influence. Recovery combines observed persistent recovery-or-censor time, excess-error area, and persistent censored fraction from the canonical paired initialization analysis. Relapse after a first qualifying entry remains a separate diagnostic because it is defined only for runs that enter the recovery band. Equal-scenario and high-severity weighting variants quantify the sensitivity of the summary to alternative deployment priorities.

2.7. HECM lookup-table sensitivity analysis

The HECM is the only tested representative whose SOC calculation depends directly on identified resistance, time-constant, and OCV lookup surfaces. Consequently, an observed HECM disturbance penalty could originate from the manipulated measurement, from local lookup-table calibration uncertainty, or from an interaction between both. The sensitivity analysis tests whether the HECM robustness profile remains stable when these lookup surfaces are varied and whether the larger disturbance penalties can still be attributed primarily to the disturbed measurements. This separates measurement robustness from parameter-calibration sensitivity and is necessary for interpreting the HECM benchmark results.

The nominal lookup surfaces are identified from the training and validation data and remain fixed during the primary cross-model benchmark. For the sensitivity test, the resistance surfaces and the time-constant surfaces are each scaled independently by $\pm 10\%$, while the OCV surfaces are shifted by $\pm 10\text{ mV}$. Each one-at-a-time lookup perturbation is crossed with the complete set of measurement and signal-integrity subcases, including the signed current-offset test, and with the paired initialization-recovery experiment. Because these lookup perturbations apply only to HECM, this analysis is evaluated separately and does not enter the cross-model robustness score.

For every disturbance, the lookup interaction is calculated as the change in scenario ΔMAE caused by the perturbed lookup relative to the corresponding

410 result with the nominal lookup. Results are first aggregated within cells and
 411 then summarized with the same equal-weight cell bootstrap as the primary
 412 benchmark. This design separates sensitivity to local HECM calibration changes
 413 from sensitivity to the disturbed measurement itself. The complete numerical
 414 protocol and the supporting figure are provided in Appendix Appendix B.

415 3. Experimental Setup

416 3.1. Dataset and data acquisition

417 The benchmark uses the MGFarm LFP 18650 dataset because its controlled
 418 design-of-experiments aging study provides long 1-Hz full-cell trajectories, re-
 419 peated capacity checkups, multiple operating conditions, and enough indepen-
 420 dent cells for a cell-disjoint comparison [27, 28, 29]. The fixed split comprises
 421 a six-cell training set, a three-cell validation set, and a six-cell independent
 422 holdout test set. Figure A.25 shows the complete SOH trajectories together
 423 with their charge-rate, discharge-rate, and depth-of-discharge aging conditions.
 424 Table A.4 provides the exact cell-to-set assignment and reports the measured
 425 duration, SOH, temperature, and load coverage. Figure A.21 summarizes the
 426 diversity within the holdout test set. Together, these materials document the
 427 cell-disjoint development boundary and the operating diversity of the evalua-
 428 tion data. LFP’s flat mid-SOC OCV and chemistry-specific aging behavior are
 429 integral to the tested problem.

430 The holdout test set is assigned to descriptive load groups using the measured
 431 95th-percentile absolute C-rate. It contains two low-load cells, three middle-
 432 load cells, and one high-load cell. The single high-load trajectory is exploratory,
 433 whereas the low- and middle-load groups support multi-cell summaries. The ex-
 434 act assignment and the corresponding load coverage are shown in Figure A.21
 435 and Table A.4. These cell-level load groups describe the operating coverage of
 436 the holdout set and are distinct from the multivariate selection of representa-
 437 tive evaluation windows. The latter uses measured SOH, temperature, C-rate,
 438 charge-throughput, and SOC-coverage descriptors as detailed in Section 3.3.

439 This separation preserves variation across cells and aging states without treat-
 440 ing every operating-state bin as an independent statistical sample.

441 3.2. Dataset ground truth and preprocessing

442 The reference preprocessing first computes the signed transferred charge as
 443 $Q_k = I_k \Delta t_k / 3600$ and the absolute transferred charge as $|Q_k|$, from which the
 444 capacity during dedicated capacity-test intervals is obtained as $C_{\text{ref}} = \sum_k |Q_k|$.
 445 The reference SOH target is then defined offline as $\text{SOH}_{\text{ref}} = C_{\text{ref}} / C_{\text{ref},0}$, where
 446 $C_{\text{ref},0}$ denotes the first measured reference capacity of the trajectory, and the
 447 resulting capacity and SOH columns are interpolated between the dedicated
 448 capacity-test anchors. The dataset SOC ground truth used in the benchmark
 449 is reconstructed offline as $\text{SOC}_{\text{ref}} = 1 + Q_c / C_{\text{ref},0}$, where the cumulative charge
 450 state Q_c is propagated from the signed charge and reset to 0 at the upper
 451 voltage anchor and to $-C_{\text{ref},0}$ at the lower voltage anchor during preprocessing.
 452 This reconstructed SOC serves as the common dataset ground truth for every
 453 reported SOC error. It is an offline evaluation quantity because its dataset-side
 454 processing and anchor logic are not fully available to an online estimator during
 455 deployment. All nominal accuracy values therefore refer to this consistently
 456 applied dataset ground truth. Paired ΔMAE compares each disturbed run with
 457 its matched baseline under the same ground truth.

458 3.3. Test trajectory and benchmark scenarios

459 All estimators are evaluated on every available cell/SOH combination in the
 460 holdout test set under the same disturbance protocol. One representative 24-
 461 h window is frozen for each available fresh, mid-life, and aged state, yielding
 462 16 windows. One test trajectory remains entirely within the fresh SOH range
 463 and therefore contributes no mid-life or aged window, as shown in Figure A.21.
 464 Every candidate window begins at the same measured full-charge condition
 465 ($\text{SOC} \geq 0.98$, $U \geq 3.58$ V). Multiple eligible windows can nevertheless exist for
 466 one cell and SOH state, and their temperature, load, charge throughput, and
 467 SOC coverage can differ substantially. Each candidate window w is therefore

468 represented by the measured operating-condition vector

$$\mathbf{z}_w = \left[\widetilde{\text{SOH}}_w \quad \bar{T}_w \quad P_{95}(T)_w \quad P_{95}(|C_{\text{rate}}|)_w \quad Q_{\text{through},w} \quad f_{\text{SOC} < 0.2,w} \right]^\top, \quad (9)$$

469 where the entries denote median SOH, mean temperature, 95th-percentile tem-
 470 perature, 95th-percentile absolute C-rate, absolute charge throughput, and the
 471 fraction of samples below 20% SOC. For every descriptor j , the candidate values
 472 are centered by their median m_j and scaled by their median absolute deviation
 473 s_j . The standard deviation is used when the median absolute deviation is zero,
 474 and an invariant descriptor receives unit scale. The selected observed window
 475 is

$$w^* = \arg \min_w \sqrt{\sum_j \left(\frac{z_{w,j} - m_j}{s_j} \right)^2}. \quad (10)$$

476 This measured-feature medoid criterion selects the actual window closest to the
 477 typical multivariate operating profile within each cell and SOH state. It avoids
 478 arbitrary or estimator-favorable window selection because neither model pre-
 479 dictions nor errors enter the criterion. Applying the same procedure separately
 480 to every holdout cell and available aging state retains diversity across cells and
 481 SOH conditions while preventing unusual operating periods from dominating
 482 the comparison.

483 The one-hour dropout uses 48 h from the same full-charge anchor to retain
 484 approximately 24 h of post-gap observation. The shared SOH LSTM receives
 485 192 h of undisturbed causal context before every scored window, matching its
 486 trained sequence length. Figure A.23 summarizes the selection and evaluation
 487 protocol.

488 The benchmark contains 20 scenario definitions: the nominal baseline, two
 489 current-noise levels, voltage and temperature noise, three signed current-gain
 490 magnitudes, one signed additive current-offset level, voltage and temperature
 491 offsets, ADC quantization, initial SOC mismatch, periodic and random sample
 492 loss, three timing-jitter levels, a one-hour burst dropout, and randomized-sign
 493 voltage spikes. Each current-gain magnitude and the 50 mA current-offset case
 494 contain matched positive and negative sign sublevels. Adverse-direction sum-

495 maries select the larger paired ΔMAE of the two signs within each cell be-
 496 fore equal-weight aggregation. Gaussian sensor-noise cases use 10 deterministic
 497 seeds. Random sample loss, jitter, and spike cases use 5 seeds. Determinis-
 498 tic cases run once. The complete traceable cross-model build comprises 7040
 499 model-window evaluations. Figure A.22 maps the scenario definitions to the ma-
 500 nipulated measurement or state and marks the selected paired reference-SOH
 501 ablations. Input disturbances act directly on measured channels, initialization
 502 interventions act on each estimator’s deployment-accessible state, and signal-
 503 integrity scenarios alter sample availability or timing while preserving causal
 504 execution.

505 *3.4. Implementation details*

506 Every estimator is executed causally, and all auxiliary quantities required
 507 by the models are rebuilt directly from disturbed inputs. This applies to Q_c ,
 508 EFC, dU/dt , dI/dt , and hourly SOH aggregates, ensuring that estimator inputs
 509 originate from the causal measurement path. The SOH LSTM forwards hidden
 510 and cell states across hourly updates after its 192-h initialization context. The
 511 primary SOC-GRU evaluation preserves its trained sequence-to-one semantics:
 512 each prediction uses the latest 2024-sample rolling window with recurrent state
 513 reset between windows. The continuous-stateful stream forms a separate de-
 514 ployment configuration whose full-trajectory accuracy and resource demand are
 515 reported alongside the rolling implementation. In freeze and dropout scenar-
 516 ios, disturbed or frozen measurements propagate consistently into derived online
 517 features. The dropout start is selected from measured current as the eligible one-
 518 hour interval with the largest absolute integrated net charge. Selection enforces
 519 at least 12 h of pre-gap observations and 24 h of post-gap observations and re-
 520 mains independent of dataset SOC and estimator output. Direct-measurement
 521 estimators use the same CV-triggered reset that defines the full-charge start,
 522 whereas the DD initial-SOC test uses an equivalent offset of online Q_c because
 523 the network exposes no explicit SOC state. The environment enforces common
 524 nominal-capacity metadata, disturbed streams, targets, masks, and metrics.

Model-specific capacity handling remains part of the estimator definitions.

4. Results

4.1. Baseline performance

Figure 4 summarizes nominal accuracy against the dataset ground truth across the 16 predefined windows on the common evaluation interval. Equal weighting within the six-cell holdout test set gives the lowest baseline MAE to DD at 0.0258 (95% CI 0.0189–0.0325), followed by HECM at 0.0337 (0.0246–0.0434), HDM at 0.0412 (0.0301–0.0515), and DM at 0.0690 (0.0499–0.0823). The RMSE ranking is identical, with values of 0.0300, 0.0388, 0.0442, and 0.0740, respectively. The ordering is consistent with the correction information available to the implementations. DM propagates charge with fixed capacity and has no continuous feedback correction. HDM reduces aging-related capacity mismatch through its SOH-dependent effective capacity. HECM additionally uses a voltage residual to correct its model-based state, while DD maps the complete multichannel history to SOC and can represent nonlinear temporal relations that are absent from the other implementations. These mechanisms explain the observed ordering within the tested data without implying that architecture alone determines accuracy in every operating domain. Results at the cell and SOH-window level reveal substantial variation across the holdout trajectories, particularly for DM, and place the cell-macro means in the context of heterogeneous operating and aging conditions. The agreement between MAE and RMSE establishes a consistent observed cell-macro ranking within the tested data. All six pairwise mean differences follow the same ranking direction, and five of the six paired 95% confidence intervals exclude zero. The effect magnitudes and uncertainty intervals therefore provide quantitative evidence of practical separation between the tested implementations. Exact sign-flip tests add a complementary check of directional consistency at cell level. With six independent cells, their Holm-adjusted p -values remain above 0.05 and are interpreted together with the effect magnitudes and confidence intervals rather




Figure 4: Baseline accuracy against the dataset SOC ground truth across the independent holdout test set. Dots show cell/SOH-window values. Bars and error bars show the equal-weight cell-macro mean and hierarchical 95% confidence interval. DD has the lowest macro MAE and RMSE, while DM has the largest error and between-cell spread.

than in isolation. Taken together, these statistics support a clear ranking for the tested implementations and operating envelope. The following scenarios complement this nominal comparison by revealing distinct strengths and weaknesses in robustness, recovery, and embedded deployment. All reported accuracy values refer to the consistently reconstructed dataset ground truth.

4.2. Current-gain sensitivity

Current-gain error is a systematic sensor error that persists over time and scales with the instantaneous current. It therefore differs both from additive offsets, which shift the measurement zero independently of signal magnitude, and from zero-mean random sensor noise. Figure 5 summarizes the matched signed current-gain sweep at $\pm 0.5\%$, $\pm 1.5\%$, and $\pm 3.0\%$. For each magnitude and holdout cell, the adverse-direction summary uses the larger MAE change from the two signs because one direction can partially compensate a pre-existing baseline error. The resulting degradation increases monotonically with gain-error magnitude for every estimator. DM and HDM show the strongest response, HECM is less sensitive, and DD is least sensitive. At 3%, the cell-macro MAE increases by 0.0108 for DM, 0.0101 for HDM, 0.0060 for HECM, and 0.0038 for DD. DM and HDM integrate the scaled current directly, so the error accumulates with




Figure 5: Current-gain sensitivity. (a) Six-cell adverse-direction ΔMAE for each matched $\pm g_I$ pair with hierarchical 95% confidence intervals. (b) Representative current before and after the +3% gain error. (c) SOC trajectories from the representative high-load test cell over three charge–discharge intervals. The test-set load assignment is documented in Figure A.21. The adverse-pair summary reports the larger penalty from each matched sign pair.

572 transferred charge. HECM is exposed through the same current-driven state
 573 propagation, but its voltage-residual correction can counter part of the result-
 574 ing state deviation. DD receives current together with voltage, temperature,
 575 cumulative charge, derivatives, and temporal context, which is consistent with
 576 its smaller observed dependence on one systematically scaled channel. The com-
 577 plete signed sweep remains available for interpreting direction-specific behavior.

578 The continuous high-load test-set analysis in Figure 6 complements the
 579 window-based benchmark and separates three mechanisms that independent
 580 24-h windows cannot resolve. The gain-error contribution varies over life be-
 581 cause its signed Coulomb-counting error interacts with charge direction and
 582 operating state. Between full charges, the penalty often grows within an inter-
 583 val, but full-charge anchoring can partially reduce or restructure it. The process
 584 is therefore neither an unrestricted monotonic drift nor a complete reset. Across




Figure 6: Current-gain-error diagnostic over the continuous high-load holdout trajectory. (a) Rolling 24-h gain-error contribution over elapsed life. (b) Mean and range over normalized intervals between consecutive full-charge events. (c) Penalty immediately before and after full-charge events. (d) Full-life baseline and adverse-gain MAE. Full-charge events reduce or restructure accumulated error but do not erase the lifetime penalty uniformly.

585 this single full-life replay, the adverse 3% gain error increases MAE by 0.0102 for
 586 DM, 0.0072 for HDM, 0.0037 for HECM, and 0.0048 for DD. HECM therefore
 587 has the smallest full-life gain-error contribution in this trajectory, while DM and
 588 HDM remain most exposed. This mechanism-focused diagnostic is not a six-cell
 589 lifecycle ranking.

590 4.3. Additive sensor offsets

591 The signed current-offset experiment shifts the measured current by ± 0.05 A
 592 and selects the larger paired ΔMAE within each cell before equal-weight aggrega-
 593 tion. DD has the smallest adverse cell-macro penalty at 0.0401 (95% CI

0.0226–0.0556), followed by HECM at 0.1763 (0.1329–0.2142), DM at 0.2467 (0.2253–0.2693), and HDM at 0.2727 (0.2546–0.2911). The adverse sign differs across models and cells, which confirms that a single favorable sign would not represent offset sensitivity reliably. The large response of the integration-based implementations follows from accumulating a persistent zero-current error throughout the evaluation window, whereas the recurrent DD implementation limits but does not eliminate this systematic shift. Among the evaluated measurement disturbances, additive current offset produces the largest global cell-macro penalty for all four representatives. The complete comparison with the remaining disturbance scenarios is shown in Figure 15.

The fixed positive voltage and temperature offsets produce a different model pattern. A $+3^\circ\text{C}$ temperature offset increases DD MAE by 0.0074 (95% CI 0.0027–0.0125), while the corresponding changes for DM, HDM, and HECM are 0.0000, 0.0002, and 0.0002. A $+0.02\text{ V}$ voltage offset increases DD by 0.0038 and HDM by 0.0005, leaves DM effectively unchanged, and changes HECM by -0.0040 in these finite windows. This pattern follows the model-specific use of the shifted channels. DM does not use temperature and uses voltage primarily for physical anchoring rather than continuous propagation. HDM receives both channels indirectly through the slowly updated SOH service. HECM responds to voltage through its observer residual, where the fixed offset can partly compensate an existing trajectory error over a finite window. DD uses absolute voltage and temperature as direct learned inputs and therefore transmits both systematic shifts into the SOC prediction. The signed HECM result must consequently be interpreted as compensation relative to its baseline rather than an improvement in the physical model. The additive current, voltage, and temperature offsets form one offset family in the robustness synthesis, with the signed current case represented by its adverse paired result.

4.4. Noise sensitivity

In contrast to the persistent gain and offset cases, the sensor-noise experiments represent random zero-mean fluctuations. They produce smaller global

624 penalties than the systematic current errors, with a response that depends on
 625 the disturbed channel and estimator structure. At $\sigma_I = 0.10$ A, HECM has
 626 the largest mean macro penalty ($\Delta\text{MAE} = 0.0038$), although its interval spans
 627 zero. DM changes by -0.0009 , while HDM and DD remain near zero. In DM
 628 and HDM, every disturbed current sample enters the cumulative charge state
 629 Q_c , but positive and negative noise increments partly cancel during integration.
 630 The integration therefore smooths rapid zero-mean fluctuations in the resulting
 631 SOC trajectory even though no individual sample is ignored. HECM processes
 632 disturbed current directly in both state propagation and the voltage-prediction
 633 path, which provides a plausible transmission mechanism for its larger mean re-
 634 sponse. DD also receives the disturbed raw current directly and, unlike DM and
 635 HDM, additionally receives dI/dt at every recurrent step. These direct channels
 636 explain the visibly larger short-term DD output variation in Figure 7(d). The
 637 rolling multichannel context distributes this response over the input sequence.
 638 At the tested level, the local transmission does not develop into a persistent
 639 global bias, so DD remains near zero in the window-wide ΔMAE . These explana-
 640 tions remain conditional on the reported uncertainty because the HECM interval
 641 includes zero. Voltage noise ($\sigma_U = 0.01$ V) produces a 0.0014 penalty for DM
 642 and 0.0013 for HDM, while HECM and DD remain below 0.0003. The small DM
 643 and HDM voltage response can arise through voltage-dependent full-charge an-
 644 choring rather than their continuous current-integration update. Temperature
 645 noise is neutral for DM and small for the SOH-aware models because temper-
 646 ature enters those pipelines through learned or parameter-dependent pathways
 647 rather than the direct charge update. Figure 7 shows the local transmission of
 648 current noise, while Figure 8 reports repeated six-cell effects with hierarchical
 649 uncertainty.

650 Taken together, the repeated six-cell results show that zero-mean sensor
 651 noise causes only small changes in window-wide accuracy for most model-
 652 channel combinations. The local current-noise analysis further demonstrates
 653 that window-wide ΔMAE should be interpreted together with short-term out-
 654 put transmission because temporal fluctuations may be only weakly represented




Figure 7: Noise sensitivity under additive current noise. (a) Baseline and noisy current in a reset-free 30 min example from the high-load test cell. (b) Matched baseline and noisy SOC trajectories in the same window. (c) Cell-macro ΔMAE with hierarchical 95% confidence intervals across the holdout test set. (d) Local transmission in the same example, measured as the 95th percentile of the change in noise-induced output deviation between consecutive samples. DD shows the largest local transmission because raw current and its derivative enter the recurrent input directly, whereas the integration in DM and HDM smooths zero-mean sample-level fluctuations. The local panels illustrate the mechanism, while panel (c) provides the test-set comparison.

655 by an aggregate error metric.

656 4.5. Initialization error and recovery

657 Initialization performance uses the paired criterion defined above. The pro-
 658 tocol applies the same -0.10 SOC-equivalent mismatch through the deployment-
 659 accessible state of each estimator. It offsets the explicit Coulomb-counting state




Figure 8: Repeated noise robustness across the six-cell holdout test set. Bars show cell-macro ΔMAE and error bars show hierarchical 95% confidence intervals with random seeds nested within cells. HECM is most affected by high current noise, while most other channel-model combinations remain close to baseline relative to the between-cell uncertainty.

660 in DM and HDM and the EKF SOC state in HECM. For DD, it applies the
 661 equivalent cumulative-charge shift $Q_c - 0.10C_{\text{nom}}$. Figure 9 reports the resulting
 662 paired trajectory differences from the common evaluation start, while recovery
 663 time remains referenced to the intervention. HECM provides the fastest per-
 664 sistent recovery, with a cell-macro mean of 1.20 h, followed by DD at 2.60 h.
 665 DD enters the recovery band temporarily in several runs but has a higher re-
 666 lapse fraction than HECM (0.78 compared with 0.17). DM and HDM generally
 667 do not recover persistently within the observation horizon, and both have a
 668 persistent-censor fraction of 0.89. HECM also has the smallest excess-error
 669 area at 0.078 SOC h, compared with 0.116 SOC h for DD. The recovery pat-
 670 tern follows the available correction mechanisms. DM and HDM propagate the
 671 imposed state offset until a physical charge anchor is reached because neither
 672 method has continuous output feedback. HECM uses the voltage residual to
 673 correct its explicit SOC state and can therefore re-anchor persistently. DD does
 674 not propagate SOC as an isolated state. Its multichannel rolling input allows the
 675 prediction to move away from the disturbed cumulative-charge feature, although
 676 the interaction of that feature with the changing recurrent context explains why




Figure 9: Response to a 10% SOC-equivalent initialization mismatch. (a) Correctly and incorrectly initialized trajectories for a representative mid-life pair from the high-load holdout trajectory. (b) Absolute paired differences over the first 6 h. The DD trajectory enters the 2% band for 5 min, leaves it again, and only later returns persistently. The vertical gray marker denotes the common evaluation start after all estimators provide valid output. (c) Observed first qualifying entry and persistent recovery-or-censor time across the holdout test set with hierarchical 95% confidence intervals. Endpoints already satisfied at the common start are left-censored.

entry into the recovery band can be followed by relapse. Together, the recovery time, relapse behavior, censoring, and excess-error area identify HECM as the strongest recovery implementation under this protocol.

4.6. Missing samples and signal integrity

Figure 10 compares the six-cell effects of periodic and random sample loss with those of timing jitter. Sample loss is substantially more consequential than timing jitter. Periodic freezing of every 50th sample increases MAE by 0.0072, 0.0065, 0.0028, and 0.0011 for DM, HDM, HECM, and DD, respectively. Random 2% loss produces corresponding increases of 0.0080, 0.0081, 0.0039, and 0.0026. By contrast, even the ± 0.9 s jitter case remains within approximately

687 6×10^{-4} SOC of baseline for every model, with confidence intervals generally
688 spanning zero. DD therefore has the smallest sample-loss penalty, while DM
689 and HDM retain their structural dependence on uninterrupted charge integra-
690 tion. A missing or frozen current sample removes or distorts charge information
691 that an integration-based estimator cannot reconstruct retrospectively. HECM
692 can reduce part of this error through voltage feedback, while DD can use the
693 surrounding multichannel sequence rather than relying on the affected sample
694 alone. Timing jitter preserves the measurements themselves, and each imple-
695 mentation reconstructs its time-dependent quantities from the disturbed sam-
696 pling intervals. This explains why timing perturbation is less consequential than
697 information loss under the tested protocol.

698 The one-hour burst dropout creates a distinct availability problem. Against
699 the duration-matched 48-h baseline, its macro ΔMAE is largest for HECM
700 (0.0239, 95% CI 0.0155–0.0366), followed by HDM (0.0209, 0.0098–0.0313) and
701 DM (0.0050, –0.0029–0.0138). DD changes by –0.0093 (–0.0288–0.0099), re-
702 flecting finite-window compensation between its baseline and frozen-input tra-
703 jectories. The DD interval includes zero, so the negative mean does not demon-
704 strate an accuracy improvement. It shows that no persistent global degradation
705 is resolved under this protocol. The interval is placed where the absolute un-
706 observed net charge is largest among starts that retain the required pre- and
707 post-gap context. Figure 11 compares the disturbed output of the represen-
708 tative high-load test cell with its duration-matched no-dropout trajectory and
709 reports the global holdout-test effect. DM and HDM retain the corrupted charge
710 balance until a physical anchor because the missing interval cannot be recon-
711 structed by integration alone. HECM receives measurements again after the
712 gap, but the flat LFP voltage relation limits how strongly its observer can infer
713 the missing charge and correct the accumulated state deviation. DD progres-
714 sively removes frozen raw and derivative samples from its finite rolling window
715 after measurements resume. Its voltage, temperature, current, and temporal
716 context can then reduce dependence on the corrupted cumulative-charge fea-
717 ture. This mechanism is consistent with its smaller persistent effect, while the




Figure 10: Six-cell signal-integrity and timing stress. Bars show cell-macro ΔMAE and hierarchical 95% confidence intervals for periodic and random sample loss and three timing-jitter levels. Sample loss clearly dominates timing jitter under the tested protocol.

finite-window confidence interval prevents a stronger claim. Dropout contributes to the robustness dimension, while the recovery dimension uses the canonical paired initialization experiment.

4.7. Spike sensitivity

Single voltage spikes demonstrate why local and global metrics must be reported together. Full-window ΔMAE remains below 4×10^{-4} SOC for every estimator and is effectively zero for DM, HDM, and DD. Event-aligned analysis, however, identifies a reproducible DD transient: its mean spike-sample penalty across cells is approximately 7.7×10^{-3} SOC, compared with near-zero direct penalties for the other representatives. Figures 12 and 13 further show that the DD response varies across cells and SOH states and then decays after the event. DM and HDM do not use voltage as a continuous SOC-correction signal, so an isolated spike away from the full-charge condition has little direct effect. HECM processes voltage through its observer residual, which filters a single inconsistent measurement through the model and covariance weighting. DD receives the impulse simultaneously through absolute voltage and dU/dt , and the disturbed values enter its rolling recurrent context. This provides a plausible mechanism for the larger but short-lived DD response. A channel-removal ablation would be




Figure 11: One-hour burst-dropout transition. (a) Estimator outputs around the shaded missing-input interval. (b) Absolute deviation from the paired duration-matched no-dropout output after measurements resume, with the 2% trajectory-difference line shown for context. (c) Six-cell global Δ MAE with hierarchical 95% confidence intervals.

required to separate the individual contributions of the absolute and derivative inputs.

4.8. ADC quantization

Figure 14 compares representative raw and quantized current and voltage signals and summarizes the resulting six-cell Δ MAE. The tested ADC steps ($\Delta I = 0.01$ A, $\Delta U = 0.005$ V, and $\Delta T = 0.5^\circ\text{C}$) preserve the local signal shape but affect the tested representatives differently. DM and HDM increase by 0.0038 and 0.0035 MAE, respectively. HECM changes by -0.0018 and DD by $+0.0002$. All four confidence intervals overlap zero. The small DM and HDM response is consistent with repeated current-rounding errors entering their cumulative charge states. HECM combines the quantized current with voltage-residual correction, while DD embeds all quantized channels in a multichannel temporal context. Their signed changes remain within the uncertainty range




Figure 12: Voltage-spike sensitivity across six cells. (a) Mean event-sample MAE penalty with hierarchical 95% confidence intervals. (b,c) Cell- and SOH-state-resolved penalties for HECM and DD, respectively. Absent cell/state combinations are left blank. The event-aligned metric exposes the local DD response that is obscured in full-window Δ MAE.

and must not be interpreted as physical improvement or exact cancellation. At the selected coarse resolution, quantization produces smaller penalties than current gain error, sample loss, and burst dropout, while the integration-based representatives retain a measurable response.

4.9. Diagnostic comparison with ground-truth SOH

The primary comparison supplies the same causally estimated LSTM-SOH trace to HDM, HECM, and DD. This estimated trace represents the health information available to the complete estimator pipelines during online operation. The diagnostic comparison replaces it with the dataset reference SOH, which is the ground-truth SOH derived from capacity checkups and interpolated between the measured anchors. Its purpose is to identify whether observed SOC errors originate predominantly within the SOC branch or are inherited through the estimated health input.

Reference SOH reduces baseline MAE by 0.0370 for HDM and 0.0166 for HECM, while DD MAE increases by 0.0049 on average. The pronounced HDM improvement is consistent with its direct use of SOH to scale effective capacity. HECM also benefits because SOH enters both capacity scaling and lookup-based parametrization. DD instead uses SOH as one learned feature among




Figure 13: Event-aligned voltage-spike response. (a) Representative SOC trajectories around the injected spike, with the dashed black line showing dataset ground-truth SOC. (b) Model-wise 95th percentile peak excess error. (c) Mean excess absolute error aligned across repeated spikes, with the DD uncertainty envelope. The local response decays rapidly and is therefore weakly represented by full-window MAE.

several interacting inputs. Replacing this single feature with its ground truth therefore does not impose a monotonic relationship between SOH accuracy and the resulting SOC error.

The same model-specific direction persists for most noise, sample-loss, and dropout cases. At 3% current gain error, reference SOH improves HDM and HECM and changes DD by +0.0091. For temperature noise, the reference-minus-LSTM difference changes from its baseline value by 0.0016 for HDM and less than 10^{-4} for HECM and DD. The incremental temperature-noise response is therefore largely independent of the SOH estimator. These findings show that SOH-input quality explains part of the baseline differences, especially for HDM and HECM, but does not account for every disturbance-specific response. The LSTM-SOH runs consequently remain the primary deployment-representative




Figure 14: ADC quantization. (a,b) Representative raw and quantized current and voltage. (c) Six-cell Δ MAE with hierarchical 95% confidence intervals and resulting quantized-signal MAE annotations.

779 results, while the ground-truth SOH runs provide the diagnostic attribution.

780 4.10. HECM lookup-table sensitivity

781 Within the tested local perturbation range, lookup-table uncertainty is not
 782 the dominant source of the observed HECM measurement-disturbance penalties.
 783 Across all 22 measurement and signal-integrity subcases, the largest absolute
 784 lookup-disturbance interaction is 0.006116 MAE. It occurs within the additive
 785 current-offset test, but its cell-bootstrap interval includes zero. This interaction
 786 is small relative to the HECM current-offset penalty of 0.1763 MAE obtained
 787 with the nominal lookup. The large offset response is therefore attributable
 788 primarily to the systematic measurement error and not to the tested variations
 789 of the parameter surfaces.

790 The lookup tables nevertheless remain relevant to baseline calibration and
 791 recovery. Five of the six lookup perturbations retain mean persistent recovery-
 792 or-censor times between 1.18 and 1.41 h, close to the nominal value of 1.20 h.
 793 Under the remaining resistance perturbation, one low-load holdout cell stays

outside the recovery band for the complete observation horizon and raises the equal-weight mean to 5.10 h. The analysis therefore answers the primary disturbance-attribution question negatively while identifying recovery as a separate calibration-sensitive boundary. Figure B.26 and Table B.5 provide the complete Appendix results.

4.11. Robustness synthesis across representatives

Figure 15 condenses 18 measurement-corruption and signal-integrity rows and shows how degraded-input behavior changes the nominal-accuracy picture. The additive current offset creates the largest global penalties for DM, HDM, and HECM, while DD remains less affected. Current-gain sensitivity increases with magnitude for every representative in the matched signed sweep. Periodic and random sample loss particularly affect DM and HDM, and the one-hour dropout produces the largest non-offset point penalty for HECM. Timing jitter remains nearly neutral. Voltage spikes reveal a distinct DD transient through event alignment despite their small full-window MAE. Figure 16 complements this scenario-level comparison by combining relative accuracy, family-balanced robustness, and recovery and by showing how alternative priority profiles change the overall assessment. Together, both figures provide the evidence required for deployment-oriented estimator selection.

Under the family-balanced definition, the robustness scores are 0.452 for DM, 0.351 for HDM, 0.360 for HECM, and 0.967 for DD. Equal weighting per evaluated scenario yields 0.549, 0.476, 0.552, and 0.793, while the highest-level-plus-offset variant yields 0.573, 0.469, 0.562, and 0.693. DD leads all three aggregations. The changing score distances and the changing order of DM and HECM demonstrate the influence of the aggregation rule, while the scenario-level results identify the disturbance mechanisms behind these relative summaries.

4.12. Embedded performance and recurrent execution modes

The STM32 experiment first verifies numerical equivalence before comparing speed or memory. Across the six nominal cell sequences, mean hardware-software MAE remains below 4×10^{-4} percentage points for every SOC core




Figure 15: Cell-macro Δ MAE across 18 measurement-corruption and signal-integrity rows on the six-cell holdout test set. Current-gain and additive current-offset rows use adverse directions from matched signed sweeps. A symmetric logarithmic color scale retains visual resolution for smaller effects despite the larger current-offset penalties. Positive values denote degradation relative to each model’s paired baseline, while small negative values arise from finite-window signed compensation. The initialization-mismatch response is evaluated separately in Figure 9.

(Figure 17). These differences are orders of magnitude below the estimator-to-dataset errors and show that the float32 C implementations reproduce their software references sufficiently closely for deployment measurements.

The isolated DM and HDM SOC updates both have sub-microsecond median latency (0.171 and 0.165 μ s), HECM requires 23.9 μ s, and continuous-state DD requires 424.9 μ s (Figure 18). The latency order follows the computation performed per update. DM and HDM require only a small number of scalar arithmetic operations. HECM additionally performs parameter lookup and observer matrix operations. DD evaluates recurrent matrix products and its MLP output head. At the nominal 1-Hz update rate, these medians occupy less than 0.043% of one sampling period. This deadline share describes the isolated estimator task and provides the execution margin available for integration with other BMS functions. The DD value in the cross-model comparison is the continuous-state deployment variant. Exact 2024-sample rolling-window semantics are evaluated separately because they recompute the full sequence at every sample.

For the cross-model embedded performance comparison, DD denotes the




Figure 16: Decision-oriented synthesis for the four tested representatives. (a) Relative min-max-normalized accuracy, family-balanced robustness across eight evaluated disturbance families, and observed persistent paired recovery. Current gain and additive current offset use adverse directions from their matched signed sweeps. (b) Composite scores under three 60/20/20 priority profiles. The profiles are interpreted together with the scenario-level evidence in Figure 15.

continuous-state deployment variant used in the latency analysis. The compiled DM, HDM, HECM, and DD firmware occupies 27.0, 27.1, 470.1, and 115.9 KiB of flash, respectively. The corresponding measured peak runtime RAM is 2.4, 2.4, 2.5, and 4.1 KiB (Figure 19). These peak values combine static storage, dynamic allocation, and call-stack demand. HECM flash occupancy is dominated by its parameter surfaces, whereas the DD parameter storage accounts for most of its flash demand. The measurements describe the isolated SOC firmware with the shared SOH trace provided externally and leave resource margin for additional BMS tasks.

Figure 20 makes the DD deployment trade-off explicit, while Figure A.24 provides the corresponding latency distributions over all six hardware cell sequences. The same pruned weights are used in every mode. Exact rolling-window inference preserves the trained software semantics but requires roughly 724 ms median inference time, or 72.4% of the one-second period, and 67.3 KiB peak runtime RAM because the complete 2024-sample sequence is recomputed




Figure 17: Numerical equivalence of STM32 and software execution over six cell sequences. Bars show mean hardware–software MAE in percentage points, dots show cells, and whiskers show the cell range. The logarithmic scale emphasizes that all implementation differences are negligible relative to model error.

for every output. Continuous-state inference evaluates only the new recurrent step and reduces latency to approximately 0.425 ms and peak RAM to 4.1 KiB while retaining similar nominal error on the six sequences. Periodically resetting that state has a comparable 4.1 KiB peak but creates much larger transient errors. The mean of the cellwise maximum dataset-SOC errors is approximately 29 percentage points, and individual cells reach approximately 39 percentage points. Each reset removes the temporal context accumulated before that point, whereas continuous forwarding preserves it. Continuous-state execution therefore provides the practical hardware configuration. Its forwarded recurrent state defines a deployment mode that is evaluated separately from the primary rolling-window robustness benchmark.

On-Device Inference-Latency Distributions Across Six Cells




Figure 18: On-device SOC-core latency distributions across six cells and three replay rounds. Panels (a–d) show per-inference distributions. Panel (e) compares minimum, 5th percentile, median, 95th percentile, and maximum on a logarithmic scale. DD denotes continuous-state inference in this cross-model latency comparison.

5. Discussion

5.1. Structural interpretation across estimator classes

The four implementations expose distinct relationships between their update mechanisms and observed disturbance responses. DM has the largest nominal




Figure 19: Embedded memory requirements of the isolated STM32 SOC implementations. (a) Compiled flash occupancy. (b) Peak runtime RAM measured during a representative holdout-test replay, combining statically allocated variable storage with maximum dynamic-memory allocation and call-stack use. DD denotes continuous-state execution, matching Figure 18. Exact rolling-window execution is compared separately in Figure 20.




Figure 20: DD inference-mode trade-off on the STM32. The panels compare maximum dataset-SOC error, median latency, compiled flash occupancy, and measured peak runtime RAM for exact rolling-window, continuous-state, and periodically reset recurrent execution. Cell points show that periodic resets, not continuous forwarding, create the dominant accuracy loss.

871 error and reacts strongly to persistent current-gain and current-offset errors as
 872 well as missing samples because its current-integration structure lacks voltage-
 873 residual correction. Its practical full-charge anchor limits accumulated drift,
 874 while the lifecycle experiment shows that the reset restructures the gain-error
 875 contribution without removing it uniformly. HDM adapts the capacity denom-

876 inator through learned SOH and thereby reduces capacity mismatch, yet it
877 retains the same charge-integration core. Its similar response to systematic cur-
878 rent errors, missing charge information, and ADC quantization follows directly
879 from this shared structure.

880 HECM combines current-driven state propagation with a voltage residual
881 and has the strongest observed recovery profile under the tested initialization
882 mapping. Additive current offset produces its largest global disturbance penalty.
883 Among the remaining cases, high current noise and burst dropout are most
884 consequential because corrupted or frozen inputs affect both state propaga-
885 tion and observer correction. The confidence interval for high current noise in-
886 cludes zero and shows substantial between-cell uncertainty. The supplementary
887 lookup analysis in Appendix Appendix B shows that the largest local lookup-
888 disturbance interaction is 0.00612 MAE, compared with a 0.1763 HECM penalty
889 from the adverse current offset itself. The disturbance response is therefore gov-
890 erned primarily by the corrupted input in this case, although absolute accuracy
891 and observer re-anchoring remain dependent on lookup calibration.

892 DD has the lowest nominal cell-macro error and the smallest adverse penal-
893 ties under both systematic current-error tests. It is also comparatively resilient
894 to isolated sample loss. Event-aligned voltage spikes expose a recurrent transient
895 that full-window MAE hides. The selected voltage and derivative features and
896 the rolling recurrent context provide a plausible transmission path for this re-
897 sponse and identify a concrete target for feature ablation and disturbance-aware
898 training.

899 The ground-truth SOH comparison demonstrates why the auxiliary state
900 estimator must be considered part of the complete SOC pipeline. Part of the
901 HDM and HECM baseline error is inherited through the estimated SOH input,
902 whereas the DD response reflects a learned interaction between SOH and the
903 remaining features. A more accurate auxiliary input therefore improves a phys-
904 ically defined capacity relation more directly than it improves a multivariate
905 learned mapping. At the same time, the limited change in several paired distur-
906 bance penalties shows that the shared SOH estimator is not the dominant cause

907 of every robustness response. This distinction prevents errors of the auxiliary
908 SOH service from being attributed incorrectly to the SOC core and supports
909 evaluation of the complete causal pipeline as the primary configuration.

910 At class level, these representative results identify transferable mechanisms
911 rather than a universal ranking. Direct-measurement methods provide trans-
912 parent and deterministic updates with minimal computational and memory de-
913 mand. Their update structure primarily propagates accumulated charge, how-
914 ever, and offers no continuous mechanism for reconstructing missing charge
915 information or correcting systematic integration errors. Reliable current cali-
916 bration, uninterrupted sampling, and physical anchoring are therefore central
917 to their robustness. Adding learned SOH information can reduce capacity mis-
918 match in a hybrid direct-measurement method, but it does not by itself add
919 voltage feedback or temporal correction to the underlying integration core.

920 Model-based ECM observers extend this structure with a physically defined
921 correction mechanism. The voltage residual allows HECM to re-anchor its in-
922 ternal SOC state and is consistent with its strong observed recovery behavior.
923 This correction does not provide the same breadth of measurement history as
924 a learned multichannel sequence representation. Its effectiveness remains con-
925 strained by the observability of the operating point, the flat OCV-SOC re-
926 lation of LFP cells, the calibrated parameter surfaces, and disturbances that
927 affect both current-driven state propagation and observer correction. The data-
928 driven representative instead combines current, voltage, temperature, derived
929 quantities, and their temporal development within one recurrent context. This
930 broader context provides a plausible mechanism for its low nominal error and
931 its comparatively small penalties under several systematic-error and sample-loss
932 scenarios. The results therefore support the value of multivariate temporal in-
933 formation without establishing that every data-driven estimator will outperform
934 every physical model.

935 Temporal context is not exclusively beneficial. The voltage-spike experiment
936 shows that a localized disturbance can propagate through selected features and
937 the recurrent state, while the periodic-reset experiment shows that removing ac-

938 cumulated context can create large transient errors. The same mechanism that
 939 allows DD to combine information across time can therefore transmit distur-
 940 bances or become unavailable when its state is handled incorrectly. Structural
 941 pruning makes this context model smaller by reducing the recurrent width, but it
 942 does not resolve these temporal failure mechanisms. Robust data-driven deploy-
 943 ment consequently requires representative training coverage, disturbance-aware
 944 evaluation, careful feature selection, and a state-handling strategy that matches
 945 the intended embedded execution mode.

946 *5.2. Limitations of nominal accuracy*

947 Nominal MAE describes average agreement with the dataset ground truth
 948 under clean inputs. The disturbance and recovery experiments add the DD
 949 voltage-spike transient, HECM re-anchoring behavior, shared-SOH feature cou-
 950 pling, and model-specific responses to random noise, systematic sensor errors,
 951 and missing information. Embedded measurements add latency, memory, and
 952 inference-state semantics. These dimensions create a deployment profile for each
 953 implementation, while the weighting analysis shows how application priorities
 954 reshape their combined scores. Estimator selection can therefore begin with
 955 the relevant fault and recovery requirements and then incorporate the available
 956 compute, memory, and state-service architecture.

957 *5.3. Implications for BMS deployment*

958 For BMS deployment, estimator classes should be selected according to
 959 the dominant sensing risks, recovery requirements, and hardware constraints.
 960 Direct-measurement methods are attractive when transparency, deterministic
 961 execution, and minimal memory are primary requirements, provided that cur-
 962 rent calibration, reliable sampling, and physical anchoring can be ensured. Hy-
 963 brid methods can add adaptive health information or learned correction to an
 964 existing physical core. Their benefit and their remaining sensitivities depend on
 965 which part of that core is modified. Model-based ECM observers provide feed-
 966 back correction and stronger re-anchoring potential at the cost of parameteriza-
 967 tion effort and larger storage requirements. Data-driven methods provide broad

968 flexibility for representing nonlinear dependencies and temporal context, but
969 require representative training data and explicit validation of feature-mediated
970 transients and recurrent-state behavior.

971 The STM32 results show that hardware resources do not constitute a fun-
972 damental deployment barrier for any of the four isolated SOC cores on the
973 selected microcontroller. Every implementation fits into the available memory
974 and completes its update within the one-second sampling interval. This re-
975 sult does not make computational constraints irrelevant, because the measured
976 resource profiles differ substantially. DM and HDM provide the smallest and
977 fastest implementations. HECM remains fast but requires 470.1 KiB of flash,
978 which is dominated by its parameter surfaces. The continuous-state DD im-
979 plementation requires more computation and runtime memory than the other
980 representatives, yet its median latency of approximately 0.425 ms and peak run-
981 time RAM of 4.1 KiB demonstrate that recurrent neural estimation is feasible
982 within the tested embedded environment.

983 The DD implementation separates two complementary deployment decisions.
984 Structured pruning reduces recurrent width before firmware export and yields
985 a dense compact network that can be executed efficiently without specialized
986 sparse kernels. Recurrent execution semantics then determine how often this
987 fixed network is evaluated and how its temporal context is retained. Exact
988 rolling-window execution requires approximately 724 ms and 67.3 KiB peak
989 runtime RAM because it repeatedly evaluates the complete input sequence.
990 Continuous-state forwarding evaluates only the new recurrent step and retains
991 similar nominal error while using a small fraction of the sampling period. Peri-
992 odic state resets preserve the low computational cost but produce large transient
993 errors because they remove accumulated temporal information. The reduction
994 from rolling-window to continuous-state cost is therefore an implementation-
995 semantic gain rather than an additional compression result. Model architec-
996 ture and execution strategy must be designed together, with the chosen state
997 semantics validated against the required estimation behavior. The firmware
998 measurements cover flash, maximum call-stack use, dynamic-memory demand,

999 and numerical equivalence. Integration of the shared SOH-LSTM, energy mea-
 1000 surement, pack communication, and concurrent BMS scheduling forms the next
 1001 system-level validation stage. Table 3 relates the measured estimator profiles to
 1002 concrete deployment priorities.

1003 Figure 16 complements the recommendation map by combining the mea-
 1004 sured accuracy, robustness, and recovery dimensions into relative priority pro-
 1005 files. The three profiles assign 60% to their named dimension and 20% to each
 1006 remaining dimension. Family-balanced, equal-scenario, and high-severity ag-
 1007 gregations produce different score distances and change the HDM-HECM or-
 1008 der. This sensitivity turns the score into a configurable decision aid, while the
 1009 scenario heatmap identifies the disturbance-specific strengths and weaknesses
 1010 behind each profile.

1011 5.4. Limitations and future extensions

1012 The conclusions are grounded in four specified implementations and the six-
 1013 cell holdout test set from one controlled LFP aging dataset. Alternative reset
 1014 logic, ECM structures, neural architectures, features, and training data can
 1015 produce different estimator profiles. The low- and middle-load strata contain
 1016 multiple cells, whereas the high-load stratum contains one exploratory trajec-
 1017 tory. One low-load test trajectory also provides no mid-life or aged window.
 1018 Figure A.21 makes these coverage limits explicit. The reported confidence in-
 1019 tervals therefore quantify the available between-cell variation within this design.
 1020 Transfer to other chemistries, cell formats, pack configurations, thermal envi-
 1021 ronments, and field duty cycles requires corresponding cell-disjoint validation.
 1022 Nominal accuracy is referenced to one consistently reconstructed offline SOC
 1023 ground truth based on Coulomb counting and voltage anchors.

1024 The primary cross-model campaign concentrates on the BMS measurement
 1025 chain, signal availability and timing, and deployment-accessible initialization
 1026 states. The systematic current-error sweep covers matched signs at fixed gain
 1027 and offset magnitudes, but not time-varying calibration drift or interacting sen-
 1028 sor errors. The separate HECM lookup analysis crosses local resistance, time-

Table 3: Decision-oriented recommendation map derived from the benchmark results. The table identifies the strongest observed representative under one dominant deployment priority and the stated protocol boundaries.

Primary deployment priority and selected representative	Main supporting benchmark evidence	Main trade-off to consider
Lowest nominal SOC error: DD	Lowest six-cell macro MAE (0.0258) and RMSE (0.0300)	Rolling-window execution is the most computationally expensive implementation
Strongest observed recovery from initialization mismatch: HECM	Lowest observed persistent-recovery-or-censor time and substantially fewer first-entry relapses than DD	Recovery is compared only after all estimators provide valid output
Lowest adverse systematic current-error penalties: DD	Matched signed tests give Δ MAE 0.0038 at 3% gain error and 0.0401 at 50 mA offset, below the other representatives	Temperature-offset and local voltage-spike sensitivity remain visible
Smallest observed periodic/random sample-loss penalty: DD	Δ MAE 0.0011/0.0026	The duration-matched dropout effect is non-positive but highly variable across cells
Lowest deterministic compute cost: DM/HDM	Sub-microsecond median SOC-core latency and approximately 27 KiB flash	Integration structure is sensitive to gain error and missing charge information
Practical fast neural deployment: DD continuous state	Approximately 0.425 ms median and 4.1 KiB peak runtime RAM	Different recurrent semantics from the primary rolling-window benchmark

constant, and OCV deviations with the complete set of tested disturbance sub-cases and paired initialization recovery. These one-at-a-time perturbations do not define probability distributions for lookup error and do not cover combined parameter deviations. Broader parameter mismatch, pack-level interactions, and supervisory BMS faults form a separate validation layer. The initialization experiment maps an SOC-equivalent offset into each estimator’s available state and compares recovery once all estimators provide valid output. The result-

ing times characterize operational re-anchoring under class-specific interventions rather than identical latent-state convergence. The continuous high-load holdout replay complements the test-set windows with a mechanism-focused lifecycle view. The hourly SOH service represents gradual capacity fade, and channel-removal experiments would further isolate the DD voltage-spike transmission path. HECM uses HPPC-derived parameter surfaces identified from the model-development sets and held fixed throughout evaluation. The sensitivity analysis therefore characterizes local disturbance interactions around this calibrated implementation rather than arbitrary lookup-table error or every possible HECM design.

The hardware profile covers the isolated SOC firmware with SOH supplied as an external causal trace. Runtime RAM is measured during a representative holdout-test replay, including complete rolling-window DD inference. The next hardware stage combines the SOC and SOH services, measures energy demand, and evaluates estimator scheduling alongside communication and supervisory BMS tasks. Broader data campaigns can then extend the disturbance protocol to interacting faults, additional operating environments, and pack-level closed-loop behavior.

6. Conclusion

This study uses representative implementations from the direct-measurement, hybrid, model-based, and data-driven SOC-estimation families to establish a common robustness and embedded-performance benchmark. Each implementation was selected to expose a characteristic update or correction mechanism under the same causal protocol. The resulting evidence connects these mechanisms with the strengths and weaknesses that become relevant when an estimator is transferred from nominal software evaluation to practical BMS operation.

For direct-measurement methods, the benchmark confirms the advantages of transparent operation, deterministic execution, and very low resource demand. It also shows why reliable current sensing, uninterrupted charge information,

1065 and physical anchoring are central to this family. Hybridization can improve an
1066 existing estimator by adding learned health information or adaptive correction,
1067 but its effect depends on the part of the underlying method that is modified.
1068 Capacity adaptation improves the treatment of ageing without automatically
1069 removing the disturbance sensitivities of a current-integration core. Model-
1070 based ECM observers add physically interpretable voltage-feedback correction
1071 and strong re-anchoring capability. These benefits are accompanied by calibra-
1072 tion effort, parameter storage, and dependence on the validity of the physical
1073 model over the intended operating range.

1074 The data-driven family shows particularly strong potential for future BMS
1075 applications. The recurrent representative combines the lowest nominal error
1076 with strong results across several disturbance families, demonstrating the value
1077 of learning nonlinear and temporal relationships from multichannel battery data.
1078 Its deployment path also shows how this capability can be made compatible with
1079 embedded execution. Structured pruning reduces the GRU width from 96 to 67
1080 units before export, while continuous-state forwarding reduces median inference
1081 latency from approximately 724 ms for exact rolling-window recomputation to
1082 0.425 ms and peak runtime RAM from 67.3 KiB to 4.1 KiB without a rele-
1083 vant loss of nominal accuracy on the evaluated sequences. Periodic reset does
1084 not provide the same behavior because removing the hidden context creates
1085 pronounced transients. Compact recurrent state estimation is therefore feasi-
1086 ble on the selected microcontroller, but its practical quality depends on both
1087 network preparation and stateful execution. Representative training coverage,
1088 disturbance-aware evaluation, and careful handling of input features and recur-
1089 rent state remain essential for exploiting this potential reliably.

1090 The broader consequence for all estimator families is that nominal MAE
1091 alone cannot establish suitability for BMS deployment. Accuracy must be evalu-
1092 ated together with disturbance-specific robustness, recovery behavior, numerical
1093 equivalence, inference latency, flash occupancy, and runtime RAM. This multidimensional benchmark does more than compare four implementations. It reveals
1094 the mechanisms that determine class-specific behavior and provides a practical
1095

1096 basis for improving a selected approach through sensing design, physical-model
1097 adaptation, disturbance-aware neural training, or embedded optimization. The
1098 same framework can be extended to complete SOC–SOH execution, concurrent
1099 BMS tasks, additional operating environments, and pack-level applications.

1100 Appendix A. Dataset Split, Holdout Coverage, and Evaluation Pro- 1101 tocol




Figure A.21: Quantitative coverage of the six-cell independent holdout test set. Panels show observed SOH, the descriptive load-class assignment, thermal envelope, and full-trajectory SOH-state coverage. The high-load group contains one exploratory trajectory.




Figure A.22: Complete 20-scenario benchmark matrix. Red markers identify the manipulated measurement, timing, availability, or initialization channel. Blue-gray markers denote repeated stochastic cases. Rose markers identify paired reference-SOH ablations. Each current-gain magnitude and the additive current-offset case contain matched positive and negative sign sublevels.




Figure A.23: Representative-window protocol. Sixteen cell/SOH windows are selected using measured-feature medoids. Every estimator is scored over the same interval once the rolling DD input context is complete. Standard windows cover 24 h after the causal SOH context, while the dropout case uses 48 h and places the one-hour gap at the eligible interval with maximum absolute net charge.




Figure A.24: DD latency distributions over the six hardware cell sequences. Exact rolling-window execution is shown separately from continuous-state and periodically reset execution because their scales and recurrent semantics differ substantially.




Figure A.25: Measured SOH trajectories across the complete LFP aging campaign. The campaign is divided into a six-cell training set, a three-cell validation set, and a six-cell independent holdout test set. The legend reports the applied charge C-rate, discharge C-rate, and depth of discharge for each trajectory. The plotted SOH values originate from repeated capacity checkups and are interpolated between checkup anchors. The exact cell-to-set assignment and measured coverage are reported in Table A.4.

Table A.4: Measured coverage and exact cell-to-set assignment for the fixed training, validation, and independent holdout test split. The 95th-percentile absolute C-rate is calculated over each complete 1-Hz trajectory. Load classes are defined only for the holdout test set. Its high-load class contains one cell and is exploratory.

Cell	Split	Duration [d]	SOH range	T range [°C]	P95 $ C_{\text{rate}} $	Test-set load
C01	Training	198.6	0.922–1.000	24.3–29.4	1.01	–
C03	Training	153.3	0.921–1.000	25.9–30.7	1.00	–
C05	Training	166.1	0.917–1.000	25.2–30.6	1.00	–
C11	Training	113.2	0.656–1.000	25.6–42.8	2.48	–
C17	Training	158.7	0.719–1.000	25.8–34.9	1.75	–
C23	Training	167.8	0.837–1.000	26.1–37.2	1.74	–
C07	Validation	167.6	0.874–1.000	25.2–32.0	1.01	–
C19	Validation	160.4	0.840–1.000	26.3–33.9	1.75	–
C21	Validation	160.8	0.824–1.000	26.2–37.6	1.75	–
C09	Test	102.6	0.682–1.000	25.0–41.1	2.51	Middle
C13	Test	44.7	0.623–1.000	24.9–42.0	2.50	Middle
C15	Test	60.7	0.627–1.000	25.1–43.6	2.53	Middle
C25	Test	127.1	0.655–1.000	25.8–37.1	1.76	Low
C27	Test	152.0	0.930–1.000	26.0–31.9	0.70	Low
C29	Test	56.0	0.604–1.000	25.5–49.4	3.00	High*

1102 Appendix B. HECM Lookup-Table Sensitivity

1103 The cross-model benchmark evaluates one fixed HECM implementation whose
1104 lookup surfaces were identified exclusively from the training and validation data
1105 and then held constant. This supplementary HECM-only analysis tests whether
1106 its disturbance response depends strongly on local lookup calibration changes.
1107 The R_i , R_1 , and R_2 surfaces are jointly scaled by -10% or $+10\%$. The τ_1 and
1108 τ_2 surfaces are likewise jointly scaled by -10% or $+10\%$, and the OCV surfaces
1109 are shifted by -10 mV or $+10$ mV. These six one-at-a-time perturbations and
1110 the nominal lookup are evaluated across 22 measurement and signal-integrity
1111 subcases plus the paired initialization intervention, all original seeds, and all 16
1112 frozen windows. The additive current-offset extension uses the same ± 0.05 A
1113 inputs as the cross-model benchmark. The combined design comprises 9184
1114 HECM executions.

1115 Main-campaign disturbances use their lookup-matched main baseline, the
1116 one-hour dropout uses a lookup-matched 48-h baseline, and initialization re-
1117 covery uses a separate paired lookup-matched baseline. For each disturbance,
1118 the lookup interaction is defined as the scenario Δ MAE under the perturbed
1119 lookup minus the corresponding Δ MAE under the nominal lookup. Results are
1120 first averaged within cells and then aggregated using the same equal-weight cell
1121 bootstrap as the primary benchmark. The analysis remains outside the cross-
1122 model robustness score because it probes one model’s local parameter sensitivity
1123 rather than a shared input disturbance.

1124 The local lookup changes shift baseline HECM MAE from 0.0314 to 0.0373,
1125 compared with 0.0337 for the nominal lookup. Across the 22 measurement and
1126 signal-integrity subcases, the largest absolute interaction is 0.006116 MAE. It
1127 occurs for the -0.05 A current offset under the $+10\%$ time-constant scaling,
1128 and its cell-bootstrap interval includes zero. The interaction is small relative to
1129 the nominal-lookup adverse current-offset penalty of 0.1763 MAE. The largest
1130 current-gain interaction remains 0.000443 MAE. Several smaller interactions
1131 have intervals that exclude zero, which identifies systematic local changes with-

out making the lookup calibration the dominant source of the observed disturbance penalties.

Initialization recovery exposes a separate boundary. Five lookup perturbations retain mean persistent recovery-or-censor times between 1.18 and 1.41 h, close to the nominal 1.20 h. Scaling all resistance surfaces by -10% instead increases the equal-weight cell mean to 5.10 h because one low-load holdout cell remains outside the recovery band through the 24-h horizon. The other five cells do not become censored. The lookup analysis therefore supports local stability of most measurement-disturbance responses, but it does not establish lookup independence of absolute accuracy or observer re-anchoring.




Figure B.26: HECM lookup sensitivity across the complete disturbance benchmark. Panel (a) reports the largest lookup interaction within each of the eight measurement and signal-integrity robustness families under local resistance, time-constant, and OCV perturbations. The additive-offset family includes current, voltage, and temperature offsets. Panel (b) reports the corresponding change in persistent initialization recovery-or-censor time. Points show equal-weight cell means. Open red points identify cell means containing a censored recovery window. Each run uses the appropriate lookup-matched main, dropout, or recovery baseline.

Table B.5: HECM lookup sensitivity across the complete disturbance benchmark. The largest robustness interaction is selected from 22 measurement and signal-integrity subcases after equal-weight cell aggregation. Bracketed counts report subcases whose 95% interaction interval excludes zero. Recovery is the persistent recovery-or-censor time from the paired initialization analysis, with the equal-weight censored-cell percentage in brackets.

Lookup condition	Baseline MAE	Max. robust $ \Delta\Delta\text{MAE} $ [CI]	Source subcase	Recovery/censor [h] [censored]
Nominal lookup	0.0337	Reference [-]	–	1.20 [0%]
Resistance -10%	0.0351	0.00191 [6/22]	Current offset +50 mA	5.10 [17%]
Resistance $+10\%$	0.0327	0.00162 [5/22]	Current offset +50 mA	1.38 [0%]
Time constants -10%	0.0343	0.00113 [6/22]	Current noise 0.10 A	1.41 [0%]
Time constants $+10\%$	0.0335	0.00612 [5/22]	Current offset –50 mA	1.22 [0%]
OCV -10 mV	0.0314	0.00260 [6/22]	Current offset +50 mA	1.18 [0%]
OCV $+10$ mV	0.0373	0.00378 [5/22]	Current offset +50 mA	1.19 [0%]

Table C.6: Six-cell baseline cell-macro statistics with hierarchical 95% confidence intervals.

Estimator	MAE [SOC]	RMSE [SOC]	P95 error [SOC]
DM	0.0690 [0.0499, 0.0823]	0.0740 [0.0532, 0.0881]	0.1088 [0.0768, 0.1307]
HDM	0.0412 [0.0301, 0.0515]	0.0442 [0.0319, 0.0557]	0.0649 [0.0449, 0.0846]
HECM	0.0337 [0.0246, 0.0434]	0.0388 [0.0286, 0.0493]	0.0644 [0.0468, 0.0821]
DD	0.0258 [0.0189, 0.0325]	0.0300 [0.0223, 0.0366]	0.0508 [0.0375, 0.0619]

Table C.7: Exact paired baseline-MAE comparisons across the six-cell holdout test set. Differences are defined as estimator B minus estimator A. Holm correction is applied across the six model-pair tests.

Estimator A	Estimator B	Mean difference [SOC]	95% CI	d_z	$p_{\text{exact}} / p_{\text{Holm}}$
DM	HDM	-0.0278 [-0.0381, -0.0171]	-1.93	0.0313 / 0.1875	
DM	HECM	-0.0352 [-0.0470, -0.0224]	-2.00	0.0313 / 0.1875	
DM	DD	-0.0431 [-0.0558, -0.0286]	-2.29	0.0313 / 0.1875	
HDM	HECM	-0.0075 [-0.0162, 0.0015]	-0.61	0.1563 / 0.1875	
HDM	DD	-0.0154 [-0.0243, -0.0084]	-1.39	0.0313 / 0.1875	
HECM	DD	-0.0079 [-0.0176, -0.0005]	-0.68	0.0938 / 0.1875	

Table C.8: Selected cell-macro scenario penalties relative to paired baseline. Current-gain and additive current-offset values use adverse directions from their matched signed sweeps.

Scenario	DM	HDM	HECM	DD
Current noise, $\sigma_I = 0.10$ A	-0.0009	+0.0000	+0.0038	-0.0000
Current gain, adverse 3%	+0.0108	+0.0101	+0.0060	+0.0038
Current offset, adverse 50 mA	+0.2467	+0.2727	+0.1763	+0.0401
Voltage offset, +0.02 V	+0.0000	+0.0005	-0.0040	+0.0038
Temperature offset, +3 °C	+0.0000	+0.0002	+0.0002	+0.0074
ADC quantization	+0.0038	+0.0035	-0.0018	+0.0002
Random missing samples, 2%	+0.0080	+0.0081	+0.0039	+0.0026
Timing jitter, ± 0.9 s	-0.0003	+0.0003	+0.0006	+0.0003
Burst dropout, 1 h	+0.0050	+0.0209	+0.0239	-0.0093
Voltage spikes, ± 0.20 V	+0.0000	-0.0003	+0.0002	-0.0001

Table C.9: Sensitivity of the illustrative normalized robustness score to the aggregation rule.

Estimator	Family balanced	Equal scenario	Highest level plus offsets
DM	0.452	0.549	0.573
HDM	0.351	0.476	0.469
HECM	0.360	0.552	0.562
DD	0.967	0.793	0.693

Table C.10: Canonical paired initialization recovery and target-hardware summary. Recovery values are equal-weight cell means over the common evaluation interval. Endpoints already satisfied when all estimators first provide valid output are recorded as conservative upper bounds.

Metric	DM	HDM	HECM	DD
First 300-s entry/censor time [h]	22.28	17.21	1.12	1.29
Persistent recovery/censor time [h]	22.28	22.28	1.20	2.60
Recovery excess-error AUC [SOC h]	1.735	1.738	0.078	0.116
First-entry censored fraction	0.89	0.67	0.00	0.00
Persistent-recovery censored fraction	0.89	0.89	0.00	0.00
Relapse fraction after first hold	0.00	0.22	0.17	0.78
Median SOC-core latency [μ s]	0.171	0.165	23.9	424.9 ^a
Flash [KiB]	27.0	27.1	470.1	115.9 ^a
Peak runtime RAM [KiB]	2.4	2.4	2.5	4.1 ^a

^aContinuous-state DD. Exact rolling-window execution requires approximately 724 ms and 67.3 KiB peak RAM.

1143 **Funding**

1144 Funding details are omitted from the blinded fallback build.

1145 **Acknowledgements**

1146 Acknowledgements are omitted from the blinded fallback build.

1147 **CRedit authorship contribution statement**

1148 The author contribution statement is provided in the non-blinded submission
1149 materials.

1150 **Declaration of competing interest**

1151 The authors declare that they have no known competing financial interests or
1152 personal relationships that could have appeared to influence the work reported
1153 in this paper.

1154 **Declaration of generative AI and AI-assisted technologies**
1155 **in the manuscript preparation process**

1156 During the preparation of this work the author(s) used ChatGPT, DeepL,
1157 and Grammarly in order to improve the linguistic quality and style of the
1158 manuscript. After using these tools, the author(s) reviewed and edited the
1159 content as needed and take(s) full responsibility for the content of the published
1160 article.

1161 **Data availability**

1162 Repository links that identify the authors are omitted from the blinded fall-
1163 back build and are provided in the author manuscript.

1164 References

- 1165 [1] P. Shrivastava, P. A. Naidu, S. Sharma, B. K. Panigrahi, A. Garg, Re-
 1166 view on technological advancement of lithium-ion battery states estima-
 1167 tion methods for electric vehicle applications, *Journal of Energy Storage* 64
 1168 (2023) 107159. doi:10.1016/j.est.2023.107159.

- 1169 [2] F. Guo, L. D. Couto, G. Mulder, K. Trad, G. Hu, O. Capron, K. Haghverdi,
 1170 A systematic review of electrochemical model-based lithium-ion battery
 1171 state estimation in battery management systems, *Journal of Energy Storage*
 1172 101 (2024) 113850. doi:10.1016/j.est.2024.113850.

- 1173 [3] J. Yao, J. Kowal, Towards a smarter battery management system: A critical
 1174 review on deep learning-based state of charge estimation of lithium-ion
 1175 batteries, *Energy and AI* 21 (2025) 100585. doi:10.1016/j.egyai.2025.
 1176 100585.

- 1177 [4] X. Li, J. Jiang, C. Zhang, L. Y. Wang, L. Zheng, Robustness of SOC
 1178 Estimation Algorithms for EV Lithium-Ion Batteries against Modeling Er-
 1179 rors and Measurement Noise, *Mathematical Problems in Engineering* 2015
 1180 (2015) 1–14. doi:10.1155/2015/719490.

- 1181 [5] F. Naseri, C. Barbu, A. Jensen, P. Setoodeh, J. R. Baena, N. Hevele,
 1182 E. Schaltz, Toward Reliable Wireless BMS: Robust Battery State Estima-
 1183 tion Under Noise and Uncertainty, *IEEE Transactions on Vehicular Tech-*
 1184 *nology* (2025) 1–12doi:10.1109/TVT.2025.3638613.

- 1185 [6] F. Berger, D. Joest, E. Barbers, K. Quade, Z. Wu, D. U. Sauer, P. Dechent,
 1186 Benchmarking battery management system algorithms - Requirements,
 1187 scenarios and validation for automotive applications, *eTransportation* 22
 1188 (2024) 100355. doi:10.1016/j.etrans.2024.100355.

- 1189 [7] Z. A. Khan, P. Shrivastava, S. M. Amrr, S. Mekhilef, A. A. Algethami,
 1190 M. Seyedmahmoudian, A. Stojcevski, A Comparative Study on Different

- Online State of Charge Estimation Algorithms for Lithium-Ion Batteries, Sustainability 14 (12) (2022) 7412. doi:10.3390/su14127412.
- [8] Y. Gao, G. L. Plett, G. Fan, X. Zhang, Enhanced state-of-charge estimation of LiFePO₄ batteries using an augmented physics-based model, Journal of Power Sources 544 (2022) 231889. doi:10.1016/j.jpowsour.2022.231889.
- [9] Y. Gao, K. Liu, C. Zhu, X. Zhang, D. Zhang, Co-Estimation of State-of-Charge and State-of-Health for Lithium-Ion Batteries Using an Enhanced Electrochemical Model, IEEE Transactions on Industrial Electronics 69 (3) (2022) 2684–2696. doi:10.1109/TIE.2021.3066946.
- [10] J. Shen, W. Ma, J. Xiong, X. Shu, Y. Zhang, Z. Chen, Y. Liu, Alternative combined co-estimation of state of charge and capacity for lithium-ion batteries in wide temperature scope, Energy 244 (2022) 123236. doi:10.1016/j.energy.2022.123236.
- [11] C. Wang, N. Cui, Z. Cui, H. Yuan, C. Zhang, Fusion estimation of lithium-ion battery state of charge and state of health considering the effect of temperature, Journal of Energy Storage 53 (2022) 105075. doi:10.1016/j.est.2022.105075.
- [12] P. Mondal, D. Bhavsar, K. Mittal, M. Mittal, Estimating State-of-Charge in Lithium-Ion Batteries Through Deep Learning Techniques: A Comparative Evaluation, IEEE Access 12 (2024) 78773–78786. doi:10.1109/ACCESS.2024.3408220.
- [13] D. P. Pau, A. Aniballi, Tiny Machine Learning Battery State-of-Charge Estimation Hardware Accelerated, Applied Sciences 14 (14) (2024) 6240. doi:10.3390/app14146240.
- [14] S. Giazitzis, M. Sakwa, S. Leva, E. Ogliari, S. Badha, F. Rosetti, A Case Study of a Tiny Machine Learning Application for Battery State-

1218 of-Charge Estimation, *Electronics* 13 (10) (2024) 1964. doi:10.3390/
1219 electronics13101964.

1220 [15] Y. Mazzi, H. Ben Sassi, A. Gaga, F. Errahimi, State of charge estima-
1221 tion of an electric vehicle's battery using tiny neural network embedded
1222 on small microcontroller units, *International Journal of Energy Research*
1223 46 (6) (2022) 8102–8119. doi:10.1002/er.7713.

1224 [16] M. A. Kamali, A. C. Caliwag, W. Lim, Novel SOH Estimation of Lithium-
1225 Ion Batteries for Real-Time Embedded Applications, *IEEE Embedded Sys-*
1226 *tems Letters* 13 (4) (2021) 206–209. doi:10.1109/LES.2021.3078443.

1227 [17] B. Gou, Y. Xu, X. Feng, An Ensemble Learning-Based Data-Driven
1228 Method for Online State-of-Health Estimation of Lithium-Ion Batteries,
1229 *IEEE Transactions on Transportation Electrification* 7 (2) (2021) 422–436.
1230 doi:10.1109/TTE.2020.3029295.

1231 [18] W. Li, N. Sengupta, P. Dechent, D. Howey, A. Annaswamy, D. U. Sauer,
1232 Online capacity estimation of lithium-ion batteries with deep long short-
1233 term memory networks, *Journal of Power Sources* 482 (2021) 228863. doi:
1234 10.1016/j.jpowsour.2020.228863.

1235 [19] F. Rzepka, P. Hematty, M. Schmitz, J. Kowal, Neural Network Architecture
1236 for Determining the Aging of Stationary Storage Systems in Smart Grids,
1237 *Energies* 16 (17) (2023) 6103. doi:10.3390/en16176103.

1238 [20] BQ79616 16-Series Battery Monitor, Balancer, and Integrated Hardware
1239 Protector (2023).

1240 [21] Design Considerations of Current Sensing With BQ769x2 Family for Bat-
1241 tery Management System (2022).

1242 [22] TMCS1108 3% Functional Isolation Hall-Effect Current Sensor With ± 100 -
1243 V Working Voltage (2024).

- 1244 [23] TLE4972 high precision coreless current sensor for automotive applications
1245 (2023).
- 1246 [24] O. Aiello, Electromagnetic Susceptibility of Battery Management Systems’
1247 ICs for Electric Vehicles: Experimental Study, *Electronics* 9 (3) (2020) 510.
1248 doi:10.3390/electronics9030510.
- 1249 [25] M. Liu, D. Geißler, S. Bian, B. Zhou, P. Lukowicz, Assessing the Impact of
1250 Sampling Irregularity in Time Series Data: Human Activity Recognition
1251 As A Case Study (Jan. 2025). arXiv:2501.15330, doi:10.48550/arXiv.
1252 2501.15330.
- 1253 [26] 10s battery pack monitoring, balancing, and comprehensive protection, 50-
1254 A discharge reference design (2024).
- 1255 [27] A. Aqachtoul, K. Karam, A. Elamrani, Y. Alidrissi, M. Najib, A. Rochd,
1256 M. Bakhouya, MGFARM: An IoT and Edge-Driven Framework for Smart
1257 and Sustainable Agricultural Practices, in: Y. Mejdoub, A. Elamri, M. Kar-
1258 douchi (Eds.), *Connected Objects, Artificial Intelligence, Telecommunica-*
1259 *tions and Electronics Engineering*, Vol. 1584, Springer Nature Switzerland,
1260 Cham, 2026, pp. 419–425. doi:10.1007/978-3-032-01536-5_64.
- 1261 [28] F. Rzepka, Design of Experiments for Cyclic Aging of LFP 18650 Cells: Im-
1262 pact of Charge/Discharge C-Rates, Depth of Discharge, and Temperature
1263 on Battery Degradation (Nov. 2025). doi:10.14279/DEPOSITONCE-24800.
- 1264 [29] K. Siebertz, D. Van Bebber, T. Hochkirchen, *Statistische Versuchspla-*
1265 *nung*, Springer Berlin Heidelberg, Berlin, Heidelberg, 2017. doi:10.1007/
1266 978-3-662-55743-3.