

Beyond Nominal Accuracy: Robustness Benchmarking of Representative Embedded Battery SOC Estimators under Measurement and Signal-Integrity Disturbances

Florian Rzepka^{a,*}, Qinnan Chen^a, Julia Kowal^a

*^aTechnical University of Berlin, Chair of Electrical Energy Storage
Technology, Einsteinufer 11, Berlin, 10587, Germany*

Abstract

Nominal accuracy alone does not reveal whether a battery State-of-Charge (SOC) estimator remains dependable when measurements are corrupted, unavailable, or initialized incorrectly. This study asks how four deployment-oriented estimator representatives compare when clean-data accuracy, degraded-input behavior, recovery, and microcontroller cost are assessed together. The concrete implementations, not entire model families, are fixed-capacity Coulomb counting with full-charge anchoring (DM), SOH-adaptive Coulomb counting (HDM), a second-order equivalent-circuit observer (HECM), and a recurrent data-driven estimator (DD). The causal benchmark uses cell-disjoint LFP aging data spanning multiple load, thermal, and state-of-health conditions and applies physically interpretable disturbances to measured inputs, timing, sample availability, and initialization. DD achieves the lowest nominal error and strongest aggregate robustness, particularly under current-gain error and isolated sample loss, but shows a local voltage-spike transient and high cost under exact rolling-window inference. HECM re-anchors fastest after initialization mismatch, yet is more vulnerable to current noise and prolonged dropout and requires the largest flash image. DM and HDM execute with minimal cost but retain the

*Corresponding author.

Email addresses: florian.rzepka@tu-berlin.de (Florian Rzepka),
qinnan.chen@campus.tu-berlin.de (Qinnan Chen), julia.kowal@tu-berlin.de (Julia Kowal)

bias and continuity sensitivity of charge integration. All isolated SOC cores run on the tested STM32, although recurrent inference semantics strongly affect latency and memory. The results show that clean-data accuracy alone cannot support deployment selection: robustness, recovery, and implementation cost reveal complementary strengths and failure modes. Conclusions are limited to the tested representatives and LFP operating envelope.

Keywords: Battery management system, State of charge, Robustness benchmark, Fault tolerance, Equivalent circuit model, Recurrent neural network, Embedded implementation, Microcontroller

1 Nomenclature

2 *Acronyms*

3	BMS	Battery management system
4	SOC	State of charge
5	SOH	State of health
6	DM	Direct-measurement model
7	HDM	Hybrid direct-measurement model
8	HECM	Hybrid equivalent-circuit model
9	DD	Data-driven neural model
10	ECM	Equivalent circuit model
11	EKF	Extended Kalman filter
12	MAE	Mean absolute error
13	RMSE	Root-mean-square error
14	ADC	Analog-to-digital converter
15	CV	Constant-voltage phase
16	OCV	Open-circuit voltage
17	GRU	Gated recurrent unit
18	LSTM	Long short-term memory
19	LFP	Lithium iron phosphate

20 *Symbols*

21	t	Time
22	$I(t)$	Current signal at time t
23	$U(t)$	Voltage signal at time t
24	$T(t)$	Temperature signal at time t
25	$\widehat{\text{SOC}}(t)$	Estimated state of charge at time t
26	$\widehat{\text{SOH}}(t)$	Estimated state of health at time t
27	$\text{SOC}_{\text{ref}}(t)$	Reference SOC target at time t
28	$\text{SOH}_{\text{ref}}(t)$	Reference SOH target at time t
29	ΔI	Current bias or offset
30	ΔU	Voltage bias or offset
31	ΔT	Temperature bias or offset
32	σ_I	Standard deviation of injected current noise
33	σ_U	Standard deviation of injected voltage noise
34	σ_T	Standard deviation of injected temperature noise
35	Q_c	Online cumulative charge state used by the SOC models
36	C_{nom}	Nominal capacity
37	C_{eff}	Effective capacity used inside the estimator
38	U_{max}	Upper voltage threshold used for full-charge detection
39	U_{tol}	Voltage tolerance used in the full-charge condition
40	U_1, U_2	Polarization voltages of the first and second RC branch
41	ΔMAE	Error increase relative to baseline

42 **1. Introduction**

43 *1.1. Motivation and practical relevance*

44 Reliable battery state estimation is a core BMS function because SOC and
45 SOH determine safety margins, usable energy, operating limits, charge con-
46 trol, and degradation-aware system management in both mobile and stationary

lithium-ion battery applications [1, 2, 3]. Practical deployment quality is not defined by clean-condition accuracy alone, because real BMS operation is exposed to sensor bias, additive noise, initialization uncertainty after restart, missing samples, irregular sampling, communication interruptions, and transient signal spikes caused by measurement-chain limitations or disturbances in the surrounding system [4, 5, 6]. The relevant engineering question is therefore not only whether an estimator achieves low MAE under nominal measurements, but also whether it converges back to a reliable state after disturbed initialization and whether it remains stable under physically plausible measurement corruption.

As summarized conceptually in Figure 1, the present benchmark treats deployment-facing estimator quality as a three-dimensional problem consisting of nominal accuracy, recovery after state inconsistency, and robustness against disturbed measurements and signal-integrity faults, embedded within broader constraints on memory and execution time. The central question is not which estimator family is universally best, but whether the nominal ranking of selected representatives remains informative when realistic failure mechanisms and microcontroller constraints are considered. This distinction is essential because one implementation can be accurate under clean measurements yet difficult to re-anchor, sensitive to a specific input fault, or too costly under its trained inference semantics.

1.2. State of the art in battery state estimation

The literature commonly groups battery-state estimators into direct-measurement or integration-based approaches, model-based approaches based on equivalent-circuit or electrochemical models, hybrid approaches that combine physical structure with learned components, and purely data-driven models based on neural sequence learning [7, 1, 2, 3]. Direct-measurement approaches, typically based on Coulomb counting, are attractive because they are simple, computationally inexpensive, and easy to deploy; however, they remain structurally sensitive to current bias, missing data, and initialization errors because they lack an internal correction mechanism [7, 4]. Model-based approaches, especially ECM-




Figure 1: Conceptual overview of deployment-oriented requirements for embedded BMS state estimation. The broader requirement space includes implementation effort, hardware constraints, and estimator performance, whereas the present study focuses on the performance branch and evaluates baseline accuracy, convergence behavior, and robustness under measurement disturbances.

77 and observer-based estimators, offer physical interpretability and voltage-based
 78 correction capability, but their practical behavior depends strongly on the qual-
 79 ity of model parametrization, operating-point observability, and measurement
 80 quality [8, 2, 4].

81 Hybrid estimators seek to combine physical state propagation with learned
 82 adaptation of capacity, parameters, or auxiliary states, which has made joint
 83 SOC/SOH estimation an active research direction [9, 10, 11]. This combination
 84 can improve baseline calibration, while it does not necessarily eliminate struc-
 85 tural sensitivity to measurement disturbances. Data-driven approaches based
 86 on recurrent or sequence models can exploit nonlinear cross-relations between
 87 current, voltage, temperature, and derived online features, and recent work has
 88 also demonstrated their embedded feasibility through TinyML-oriented imple-
 89 mentations, pruning, and quantization [12, 13, 14, 15, 3].

90 1.3. Performance and embedded deployment gap

91 A large share of the SOC estimation literature still emphasizes clean-data ac-
92 curacy, average error metrics, or computational efficiency, while comparatively
93 fewer studies analyze how estimators behave under realistic measurement faults
94 and timing corruption [1, 13, 14, 15]. Existing robustness studies often remain
95 confined to a single model family, for example observer-based robustness against
96 modeling and measurement errors [4] or robust estimation concepts for specific
97 BMS architectures under noise and uncertainty [5], whereas cross-family com-
98 parisons under the same disturbance protocol and under comparable embedded
99 constraints remain limited. Benchmarking studies in the broader BMS con-
100 text underline the importance of reproducible scenario design and deployment-
101 relevant validation, but they do not by themselves close the gap between clean-
102 data estimator comparison and measurement-level performance benchmarking
103 across fundamentally different estimator classes [6]. This gap is especially rele-
104 vant in embedded BMS design, where estimator selection must jointly consider
105 nominal accuracy, convergence, robustness under disturbed measurements, and
106 computational constraints.

107 1.4. Contribution and study objectives

108 This work addresses the gap through three linked questions: whether clean-
109 data accuracy predicts performance under measurement and signal-integrity
110 disturbances; which structural mechanisms govern degradation and recovery;
111 and how the resulting behavior trades against embedded execution cost. To
112 answer them, four deployment-oriented representatives are executed under one
113 causal protocol on cell-disjoint LFP aging data covering multiple load and SOH
114 states. The campaign combines physically interpretable measurement, timing,
115 availability, quantization, and initialization disturbances with a common recov-
116 ery criterion, cell-level uncertainty analysis, a paired SOH-input ablation, and
117 STM32 measurements of the isolated SOC cores.

118 The resulting evidence supports comparisons among these four implementa-
119 tions under the declared protocol. It does not establish a universal ranking of

Table 1: Estimator-class envelope and scope of the four benchmark representatives. Results apply to the selected implementations in the last column, not automatically to every method listed in the broader class envelope.

Class	Minimum defining mechanism	Broader class can include	Representative used in this study
Direct measurement	Causal integration of measured current using an assumed capacity	Coulombic efficiency, capacity adaptation, OCV or end-point anchoring, plausibility logic, and reset thresholds	Fixed-capacity Coulomb counting with a sustained full-charge CV reset; no voltage-residual correction
Hybrid direct measurement	Integration core augmented by an estimated or learned state/parameter	Learned capacity or efficiency, residual correction, multi-sensor fusion, and adaptive anchoring	Same Coulomb-counting and reset logic as DM, with effective capacity supplied by the shared causal LSTM-SOH estimate
ECM observer	Coulomb-counted SOC propagation plus a voltage model and feedback correction	Higher RC order, hysteresis, thermal states, SOC/SOH/temperature maps, and different nonlinear observers	Second-order RC model with EKF correction, charge/discharge SOC-SOH lookup surfaces, and shared causal SOH input
Data driven	Learned mapping from measured signals or their history to SOC	Feed-forward, convolutional, recurrent, attention-based, physics-informed, raw-signal, or engineered-feature models	Structurally pruned rolling-window GRU using voltage, current, temperature, SOH, Q_c , derivatives, and Δt

120 direct-measurement, hybrid, observer-based, or data-driven families, whose be-
121 havior can change with anchoring logic, parameterization, architecture, features,
122 training data, and deployment semantics.

123 2. Methodology

124 2.1. Estimator representatives and class boundaries

125 Figure 2 summarizes the causal benchmark chain, from measured signals
126 and online feature reconstruction to disturbance injection and global and local
127 evaluation. The four estimators are selected to expose distinct correction mech-
128 anisms, not to span every implementation possible within each class. Because
129 class boundaries overlap, each representative is assigned according to its dom-
130 inant SOC update and correction mechanism. Table 1 separates this defining
131 principle from optional extensions and from the exact implementation evaluated
132 here. There is no meaningful universal “maximum” model within a class: adding
133 learned corrections, higher-order dynamics, richer sensing, or more complex se-
134 quence architectures progressively changes both capability and deployment cost.

135 All SOH-dependent representatives, namely HDM, HECM, and DD, receive
136 the same causal LSTM-SOH trace so that differences arise from how the SOC
137 branch consumes that context. DM is the low-complexity integration reference.
138 HDM isolates the effect of learned capacity adaptation without adding volt-
139 age feedback. HECM adds model-based voltage-residual correction to current-




Figure 2: Methodology overview for the causal benchmark chain. Measured battery signals are converted into online auxiliary features, processed by representative estimator implementations, exposed to measurement-only disturbances, and finally evaluated through global and local performance metrics.

driven state propagation. DD represents compact temporal learning; the GRU was selected because SOC and disturbance recovery depend on recent signal history while its gated single-state structure is lighter than a comparable LSTM. Because DD also consumes the engineered Q_c feature, it is a data-driven sequence estimator with physics-informed inputs rather than a raw-signal-only neural model.

2.2. SOC estimation methods

2.2.1. Direct-measurement model (DM)

The direct-measurement SOC formulation used in DM, HDM, and as an explicit auxiliary feature in DD follows the same online cumulative-charge state, denoted here uniformly as Q_c . It is updated causally from the measured current stream as

$$Q_c(k) = Q_c(k-1) + \frac{I_k \Delta t_k}{3600 \text{ s h}^{-1}}, \quad (1)$$

where I_k is given in ampere and Δt_k in seconds, so that the factor 3600 s h^{-1} converts the increment to ampere-hours. The sign convention is fixed by the benchmark implementation and the state is reset to $Q_c = 0$ after a sustained constant-voltage full-charge event. In the DM, this same charge state is mapped directly to SOC through the fixed nominal capacity

$$\widehat{\text{SOC}}_{\text{DM}}(k) = \text{clip}\left(1 + \frac{Q_c(k)}{C_{\text{nom}}}, 0, 1\right), \quad (2)$$

148 *2.2.2. Hybrid direct-measurement model (HDM)*

The HDM preserves exactly the same Coulomb-counting core and differs only in the capacity term. The SOH estimate rescales the usable capacity as

$$C_{\text{eff}}(k) = C_{\text{nom}} \cdot \widehat{\text{SOH}}(k), \quad (3)$$

which yields

$$\widehat{\text{SOC}}_{\text{HDM}}(k) = \text{clip}\left(1 + \frac{Q_c(k)}{C_{\text{eff}}(k)}, 0, 1\right). \quad (4)$$

149 This formulation makes the relation between DM and HDM explicit: both use
 150 the same online charge state Q_c , while only the denominator changes through
 151 online SOH support. In both direct-measurement variants, a sustained near-
 152 maximum voltage condition is interpreted as a constant-voltage full-charge event;
 153 if the voltage remains above $U_{\text{max}} - U_{\text{tol}}$, where U_{tol} denotes the admissible volt-
 154 age margin below the upper threshold U_{max} , for the configured CV duration, the
 155 internal charge state is reset so that $Q_c = 0$ and therefore $\widehat{\text{SOC}} = 1$, which re-
 156 flects the practical assumption that a completed CV phase provides a physically
 157 anchored full-charge reference.

158 *2.2.3. Hybrid equivalent-circuit model (HECM)*

The hybrid ECM model is implemented as a second-order RC model with the state vector

$$\mathbf{x}_k = \begin{bmatrix} \widehat{\text{SOC}}_k & U_{1,k} & U_{2,k} \end{bmatrix}^\top, \quad (5)$$

where U_1 and U_2 denote the polarization voltages of the two RC branches. The prediction step propagates the state according to

$$\mathbf{x}_k^- = \mathbf{A}_k \mathbf{x}_{k-1} + \mathbf{b}_k I_{k-1}, \quad (6)$$

with

$$\mathbf{A}_k = \text{diag}\left(1, e^{-\Delta t_k/\tau_1}, e^{-\Delta t_k/\tau_2}\right), \quad \mathbf{b}_k = \begin{bmatrix} \eta \Delta t_k / C_b \\ R_1 (1 - e^{-\Delta t_k/\tau_1}) \\ R_2 (1 - e^{-\Delta t_k/\tau_2}) \end{bmatrix}, \quad (7)$$

where $C_b = C_0 \widehat{\text{SOH}}$ is the SOH-scaled available capacity. The corrected terminal-voltage estimate is written as

$$\widehat{U}_k = \text{OCV}(\widehat{\text{SOC}}_k, \widehat{\text{SOH}}_k, I_k) + U_{1,k} + U_{2,k} + R_i I_k, \quad (8)$$

and the Kalman update uses the residual between $\widehat{U}_k$ and the measured terminal voltage to correct the internal state. In the present implementation, the state propagation is driven by the previous current sample I_{k-1} , whereas the measurement equation uses the current sample I_k available at the correction step. The ECM parameters are not constant. The implementation interpolates R_i , R_1 , R_2 , τ_1 , τ_2 , OCV, and dOCV/dSOC from a preidentified lookup table over SOC and SOH, with separate charge and discharge parameter surfaces. The HECM is implemented as a two-RC observer with SOH-dependent parameter scheduling.

2.2.4. Data-driven SOC model (DD)

The data-driven SOC model is a GRU-based recurrent estimator with the feature vector $\{U \text{ [V]}, I \text{ [A]}, T \text{ [}^\circ\text{C]}, \text{SOH}, Q_c, dU/dt, dI/dt, \Delta t\}$. The benchmarked structurally pruned and fine-tuned checkpoint consists of a single-layer GRU with hidden size 67 followed by an MLP head with one hidden layer of size 96 and a sigmoid output layer. During optimization, the recurrent core is exposed to causal sequence segments of length $L_{\text{SOC}} = 2024$ samples. Within this feature set, Q_c is exactly the same online cumulative-charge state defined in (1). It is included explicitly so that the neural estimator receives the same charge-throughput information that also drives DM and HDM. The derivative channels are computed online as the finite differences $dI/dt(k) = (I_k - I_{k-1})/\Delta t_k$ and $dU/dt(k) = (U_k - U_{k-1})/\Delta t_k$. They provide local transient information that is not captured by the absolute channels alone. The additional feature Δt_k allows the network to distinguish amplitude changes from sampling-interval changes under irregular timing. During the primary benchmark, each output is calculated from the latest causal 2024-sample window and recurrent state is reset between windows, matching training and validation. No future information is

used. Continuous hidden-state forwarding is evaluated only as a separate deployment ablation because it changes the learned inference semantics.

2.3. SOH estimation method

The shared SOH estimator used in the current benchmark is an LSTM-based sequence-to-sequence model that operates on hourly aggregated online features derived from $\{U \text{ [V]}, I \text{ [A]}, T \text{ [}^\circ\text{C]}, \text{EFC}, Q_c\}$, where the aggregation uses the statistics mean, standard deviation, minimum, and maximum for each feature [16, 17, 18, 19]. This transforms the five base channels into a 20-dimensional hourly feature vector. Architecturally, the SOH network first projects the aggregated input through a two-layer feature embedding block with size 128, layer normalization, GELU activations, and dropout, then processes the sequence with a two-layer unidirectional LSTM of hidden size 112, and finally maps the recurrent output through three residual MLP blocks and a prediction head with hidden size 160 to obtain the SOH estimate for each hourly step.

During optimization, the SOH estimator is exposed to causal hourly sequence segments, and benchmark execution processes one hourly aggregate at a time with causal hidden- and cell-state forwarding. This allows the estimator to accumulate degradation context without access to future information; the first hourly output is anchored with the configured initial SOH value before subsequent hourly predictions are taken directly from the recurrent model output. The SOH prediction is updated at hourly cadence and then held piecewise constant between update points, which enforces an online deployment constraint and avoids sample-wise access to future information while still allowing the faster SOC estimators to use the latest SOH context. This cadence matches the slowly varying capacity fade observed in the dataset but cannot represent an abrupt capacity-loss event, which is absent from the measurements and outside the study claims. In the HDM, the predicted SOH enters the SOC computation through the effective capacity relation $C_{\text{eff}} = C_{\text{nom}} \cdot \widehat{\text{SOH}}$, so the method remains structurally close to direct measurement while replacing the fixed-capacity assumption. In the HECM, the predicted SOH affects the capacity scaling and

215 the lookup-based ECM parametrization, while in the DD model it is provided
216 as an explicit SOC input feature that contextualizes the current measurement
217 trajectory.

218 *2.4. Model and target-hardware preparation*

219 The software benchmark uses frozen floating-point checkpoints and scalers
220 selected before holdout evaluation. Neural training, validation, hyperparameter
221 selection, scaler fitting, pruning, and checkpoint selection use only the declared
222 training and validation cells; no holdout cell enters these steps. The imported
223 HECM lookup surfaces were preidentified from separate HPPC experiments and
224 held fixed before this benchmark; no scored trajectory or benchmark result was
225 used to fit or retune them. The compact SOC-GRU has hidden size 67 and
226 the shared SOH-LSTM has hidden size 112. Configuration, checkpoint, and
227 scaler hashes are stored with the campaign manifest, and the runner rejects
228 cells outside the declared holdout split unless an explicit diagnostic override is
229 supplied.

230 The target-hardware experiment evaluates the isolated SOC cores on a NUCLEO-
231 H753ZI board with a 480 MHz STM32H753 Cortex-M7. To separate SOC im-
232 plementation cost from the shared SOH service, the board receives the precom-
233 puted causal SOH trace used by the software reference; the SOH-LSTM itself
234 is not executed on the microcontroller. The hardware results therefore measure
235 numerical equivalence, latency, and memory of the SOC implementations rather
236 than end-to-end SOC-SOH system latency. Device time is measured per valid
237 inference with the data-watchpoint-and-trace cycle counter. Each cell sequence
238 is replayed three times, firmware section sizes are extracted from the release
239 ELF images, and predictions are compared with both the software implemen-
240 tation and dataset SOC. Reported RAM is static ELF allocation including the
241 linker-reserved heap/stack block, not measured peak stack or dynamic activa-
242 tion memory.




Figure 3: Taxonomy of the disturbance scenarios used in the performance benchmark, grouped into input disturbances, initialization errors, and signal-integrity faults.

2.5. Performance benchmark design

The benchmark follows a strict measurement-only principle: only quantities that a real BMS measures or derives online from measured data are manipulated, namely current, voltage, temperature, sample availability, and sampling time, whereas artificial parameter mismatch and synthetic hidden-state corruption are deliberately excluded [6]. Figure 3 groups the benchmark into input disturbances, initialization errors, and signal-integrity faults, while Table 2 defines the exact disturbance implementations used in the study.

The disturbance magnitudes are not intended as one-to-one copies of a single field specification. Instead, they were selected from three justification classes: deployment-level sensing limitations, adverse measurement corruption, and deliberate signal-integrity stress tests [20, 21, 22, 23, 24, 25]. The main benchmark matrix follows the scenario families shown in Figure 3 and provides the basis of the main result figures and tables. ADC quantization is part of the same scenario set and is reported in Figure 16. Overall, the benchmark is designed as a controlled and reproducible scenario study that exposes estimator-specific failure modes under a common disturbance protocol rather than emulating the complete system stack of a production BMS.

Table 2: Disturbance configurations used in the benchmark and their practical rationale. Exact numeric values are fixed benchmark settings; the cited sources motivate disturbance type and magnitude range.

Scenario	Benchmark setup	Practical origin / rationale
ADC quantization	$\Delta I = 0.01 \text{ A}$; $\Delta U = 0.005 \text{ V}$; $\Delta T = 0.5 \text{ }^\circ\text{C}$	ADC steps; sensor scaling; rounding; coarse deployment-level discretization [20, 21, 26]
Current gain error	Signed gain sweep: $I_{\text{meas}} = I(1 + b)$ with $b = \pm 0.5\%$, $\pm 1.5\%$, and $\pm 3.0\%$	Sensor gain error; miscalibration; drift; upper bound as stress case [21, 22, 23]
Current noise	Gaussian; $\sigma_I = 0.02 \text{ A}$ and 0.10 A	Sensing-chain noise; EMI pickup; low and high sensor-noise levels [4, 5]
Voltage noise	Gaussian; $\sigma_U = 0.01 \text{ V}$	Wiring noise; front-end noise; poor filtering; EMI [4, 20, 24]
Temperature noise	Gaussian; $\sigma_T = 1.0 \text{ }^\circ\text{C}$	Sensor tolerance; conversion noise; thermal gradients; sensor-to-cell mismatch [26, 11]
Initial SOC error	Initial mismatch: -0.10 SOC ; DD via perturbed online Q_c	Restart; storage; missing recent history; weak LFP OCV anchoring [4, 8]
Missing samples	Every 50th sample frozen; separate random 2% loss	Packet loss; blocked updates; logger dropouts; periodic and random data freeze [5, 25]
Irregular sampling	Uniform jitter: $\pm 0.1 \text{ s}$, $\pm 0.5 \text{ s}$, and $\pm 0.9 \text{ s}$ around nominal 1 s grid	Scheduling jitter; bus latency; asynchronous acquisition [6, 5, 25]
Burst dropout	One central frozen gap: 3600 s	Communication interruption; frozen task; stuck data stream; severe recovery case [6, 5]
Voltage spikes	Single-sample spikes: $\pm 0.20 \text{ V}$ every 1000 samples (1000 s)	EMI; switching disturbance; corrupted voltage samples [24]

2.6. Metrics and evaluation criteria

For deployment-oriented interpretation, the benchmark results are synthesized along three performance dimensions: nominal accuracy under baseline conditions, convergence behavior after initialization mismatch or interrupted signal continuity, and robustness under sustained or transient measurement corruption. Disturbance-related global performance is quantified through MAE, RMSE, bias, maximum error, and the 95th percentile absolute error, because these metrics provide a compact view of average performance, systematic deviation, and tail behavior across the selected evaluation windows. Local disturbance performance is analyzed through scenario-specific metrics such as jump

counts, output variance, absolute-error variance, and excess rolling MAE, because several disturbances do not primarily manifest as window-wide MAE collapse but as short-term output fluctuations, delayed drift, or slow re-synchronization. Whenever a disturbance mask exists, the analysis further separates pre-disturbance, disturbed, and post-disturbance behavior in order to quantify disturbance-window error growth, recovery time, and residual post-disturbance error.

ADC quantization, current bias, voltage noise, temperature noise, missing samples, and irregular sampling are interpreted primarily through global metrics, whereas burst dropout, voltage spikes, current-noise detail, and initialization mismatch additionally use dedicated local evidence because their most relevant effects are episodic or recovery-related. Recovery comparisons use one estimator-independent criterion: absolute SOC error below 0.02 for a continuous 300 s, evaluated up to a common 24 h censoring horizon. Runs that do not recover are retained as censored observations rather than silently omitted. The comparison reports recovery-or-censoring time, the censored fraction, initial post-event error, MAE after 1 h and 6 h, and excess-error area under the curve. The requested -0.10 initialization mismatch is mapped into each estimator’s available online state: the coulomb-counting SOC initialization for DM and HDM, the EKF SOC state for HECM, and the capacity-equivalent Q_c offset for DD. Because the DD mapping passes through the recurrent network rather than directly fixing its output, every realized initial output error is reported before recovery is compared. The resulting ranking compares deployment-accessible initialization interventions, not mathematically identical latent-state perturbations.

The independent statistical unit is the holdout cell, not an individual time sample or selected window. Predeclared SOH windows are first averaged within each cell and random seed; the six cells then receive equal macro weight. Hierarchical bootstrap resampling with seeds nested inside cells and 10,000 repetitions provides 95% confidence intervals. Scenario effects are evaluated as paired, cell-matched differences from baseline, and estimator pairs are compared using exact two-sided sign-flip tests, raw mean and median differences, paired standardized

effect sizes, and Holm-adjusted p -values. Because only six independent cells are available, effect magnitudes and confidence intervals are treated as primary evidence and p -values as supporting evidence.

3. Experimental Setup

3.1. Dataset and data acquisition

The benchmark uses the MGFarm LFP 18650 dataset because its controlled design-of-experiments aging study provides long 1-Hz full-cell trajectories, repeated capacity checkups, multiple operating conditions, and enough independent cells for a cell-disjoint comparison [27, 28, 29]. Model development and final testing use a fixed split. SOC and SOH training use C01, C03, C05, C11, C17, and C23; validation uses C07, C19, and C21; and the independent holdout benchmark uses C09, C13, C15, C25, C27, and C29. The six test cells span distinct aging, thermal, duration, and load envelopes, as summarized in Figure A.21. LFP’s flat mid-SOC OCV and chemistry-specific aging behavior are integral to the tested problem; conclusions therefore do not transfer directly to NMC/NCA or pack-level operation.

The holdout cells are assigned to descriptive load groups before model results are inspected, using visible separation in the measured 95th-percentile absolute C-rate: C25 and C27 form the low-load group, C09, C13, and C15 the middle-load group, and C29 the high-load case. Because the high-load group contains one cell, its group-by-SOH results are explicitly exploratory and are not interpreted as population-level high-load inference. Within each trajectory, performance is additionally stratified by reference SOH (fresh: $\text{SOH} \geq 0.90$; mid-life: $0.80 \leq \text{SOH} < 0.90$; aged: $\text{SOH} < 0.80$), instantaneous absolute C-rate (< 0.5 , $0.5\text{--}1.5$, and ≥ 1.5), temperature ($\leq 30^\circ\text{C}$, $30\text{--}35^\circ\text{C}$, and $> 35^\circ\text{C}$), and SOC (< 0.20 , $0.20\text{--}0.80$, and > 0.80). Absent states are reported as not available and are never imputed.

3.2. Reference targets and preprocessing

The reference preprocessing first computes the signed transferred charge as $Q_k = I_k \Delta t_k / 3600$ and the absolute transferred charge as $|Q_k|$, from which the capacity during dedicated capacity-test intervals is obtained as $C_{\text{ref}} = \sum_k |Q_k|$. The reference SOH target is then defined offline as $\text{SOH}_{\text{ref}} = C_{\text{ref}} / C_{\text{ref},0}$, where $C_{\text{ref},0}$ denotes the first measured reference capacity of the trajectory, and the resulting capacity and SOH columns are interpolated between the dedicated capacity-test anchors. The reference SOC target used in the benchmark corresponds to the offline Coulomb-counting target $\text{SOC}_{\text{ref}} = 1 + Q_c / C_{\text{ref},0}$, where the cumulative charge state Q_c is propagated from the signed charge and reset to 0 at the upper voltage anchor and to $-C_{\text{ref},0}$ at the lower voltage anchor during preprocessing. These reference targets are intentionally treated as offline evaluation quantities, because they rely on dataset-side processing and anchor logic that are not fully available to an online estimator during deployment.

3.3. Test trajectory and benchmark scenarios

All estimators are evaluated on every available holdout-cell/SOH combination under the same disturbance protocol. One representative 24-h window is frozen for each available fresh, mid-life, and aged state, yielding 16 windows; C27 contributes only a fresh window because its measured SOH does not fall below 0.90. Windows start at a measured full-charge anchor ($\text{SOC} \geq 0.98$, $U \geq 3.58 \text{ V}$) and are selected as measured-feature medoids without using estimator outputs or errors. The one-hour dropout uses 48 h from the same anchor to retain approximately 24 h of post-gap observation. The shared SOH LSTM receives 192 h of undisturbed causal context before every scored window, matching its trained sequence length. Figure A.23 summarizes this protocol.

The 19 predeclared cases include the nominal baseline, two current-noise levels, voltage and temperature noise, three current-bias levels, voltage and temperature offsets, ADC quantization, initial SOC mismatch, periodic and random sample loss, three timing-jitter levels, a one-hour burst dropout, and

randomized-sign voltage spikes. Gaussian sensor-noise cases use 10 deterministic seeds; random sample loss, jitter, and spike cases use 5 seeds; deterministic cases run once. Figure A.22 maps each case to the manipulated measurement or state and marks the selected paired reference-SOH ablations. Input disturbances are applied directly to measured channels, initialization errors only where structurally equivalent, and signal-integrity scenarios to sample availability or timing while preserving causal estimator execution.

3.4. Implementation details

Every estimator is executed causally, and all auxiliary quantities required by the models are rebuilt directly from disturbed inputs. This applies to Q_c , EFC, dU/dt , dI/dt , and hourly SOH aggregates, so no hidden dataset-side feature leaks into the estimator input. The SOH LSTM forwards hidden and cell states across hourly updates after its 192-h initialization context. The primary SOC-GRU evaluation preserves its trained sequence-to-one semantics: each prediction uses the latest 2024-sample rolling window with recurrent state reset between windows. A faster continuous-stateful stream is treated only as a deployment ablation because it changed full-trajectory DD accuracy and therefore cannot silently replace the trained benchmark semantics. In freeze and dropout scenarios, disturbed or frozen measurements propagate consistently into derived online features. Direct-measurement estimators use the same CV-triggered reset that defines the full-charge start, whereas the DD initial-SOC test uses an equivalent offset of online Q_c because the network exposes no explicit SOC state. The environment enforces identical capacity assumptions, disturbed streams, targets, and metrics across all estimators.

4. Results

4.1. Baseline performance

Figure 4 summarizes nominal performance across the 16 frozen windows. Equal weighting of the six holdout cells yields a different ranking from the




Figure 4: Baseline performance across six independent holdout cells. Dots show cell/SOH-window values; bars and error bars show the equal-weight cell-macro mean and hierarchical 95% confidence interval. DD has the lowest macro MAE and RMSE, while DM has the largest error and between-cell spread.

previous single-trajectory analysis: DD has the lowest baseline MAE at 0.0258 (95% CI 0.0189–0.0323), followed by HECM at 0.0335 (0.0246–0.0429), HDM at 0.0405 (0.0297–0.0506), and DM at 0.0678 (0.0491–0.0809). The RMSE ranking is identical, with values of 0.0300, 0.0389, 0.0438, and 0.0733, respectively. Cell-level points reveal substantial between-cell variation, particularly for DM, so the bars must be interpreted as cell-macro summaries rather than performance on a homogeneous trajectory.

4.2. Current-bias sensitivity

Figure 5 summarizes the adverse response to a signed current-gain sweep at $\pm 0.5\%$, $\pm 1.5\%$, and $\pm 3.0\%$. For each magnitude and holdout cell, the larger MAE change from the two signs is reported because the opposite sign can partially compensate a pre-existing baseline error. The resulting worst-case degradation increases monotonically with gain-error magnitude for every estimator. DM and HDM show the strongest response, HECM is less sensitive, and DD is least sensitive. At 3%, the cell-macro MAE increases by 0.0107 for DM, 0.0100 for HDM, 0.0059 for HECM, and 0.0038 for DD.

The continuous C29 lifecycle analysis in Figure 6 separates three mechanisms that independent 24-h windows cannot resolve. The bias contribution varies over




Figure 5: Current-gain sensitivity. (a) Six-cell adverse-direction ΔMAE for each matched $\pm b$ pair; error bars are hierarchical 95% confidence intervals. (b) Representative current before and after the +3% gain error. (c) C29 SOC trajectories over three charge–discharge intervals. The adverse-pair summary prevents signed compensation of pre-existing baseline bias from being misread as improved robustness.

life because its signed Coulomb-counting error interacts with charge direction
and operating state. Between full charges, the penalty generally grows within
an interval, but full-charge anchoring can partially reduce or restructure it. The
process is therefore neither an unrestricted monotonic drift nor a complete reset.
Across the full replay, DD retains the smallest adverse penalty, whereas DM and
HDM remain most exposed to accumulated current error.

4.3. Noise sensitivity

Repeated sensor noise is weaker than persistent current gain error, but the
response depends on channel and estimator structure. At $\sigma_I = 0.10$ A, HECM
has the largest macro penalty ($\Delta\text{MAE} = 0.0056$), whereas DM changes by
 -0.0009 and HDM and DD remain near zero. Small negative changes are paired
finite-window effects, not evidence that noise improves an estimator in general.




Figure 6: Current-bias accumulation over the continuous C29 life replay. (a) Rolling 24-h bias contribution over elapsed life. (b) Mean and range over normalized intervals between consecutive full-charge events. (c) Penalty immediately before and after full-charge events. (d) Full-life baseline and adverse-bias MAE. Full-charge events reduce or restructure accumulated error but do not erase the lifetime penalty uniformly.

416 Voltage noise ($\sigma_U = 0.01$ V) produces a 0.0014 penalty for DM and 0.0013
 417 for HDM, while HECM and DD remain below 0.0003. Temperature noise is
 418 neutral for DM and small for the SOH-aware models. Figure 7 shows the local
 419 transmission of current noise, while Figure 8 reports repeated six-cell effects
 420 with hierarchical uncertainty.

4.4. Initialization error and recovery

422 Initialization performance is evaluated using one model-independent target
 423 criterion: the shifted output must remain within 0.02 SOC of the paired correct-
 424 initialization output for at least 300 s, with censoring at 24 h. The same re-




Figure 7: Noise sensitivity under additive current noise. The figure combines an applied current-noise example, a matched SOC comparison in the same 30 min window, the global MAE penalty, and the local output-variability sensitivity across noise levels.

425 requested -0.10 SOC-equivalent mismatch is mapped to the explicit DM/HDM
 426 Coulomb-counting state, the HECM EKF SOC state, and a Q_c offset of $-0.10C_{\text{nom}}$
 427 for DD. Figure 9 reports the realized shift rather than assuming that these in-
 428 ternal mappings create identical outputs. Across cells, HECM recovers fastest
 429 at 1.03 h on average, followed by DD at 5.57 h; DM and HDM both average
 430 22.28 h and are frequently close to the censor horizon. Thus, observer correction
 431 is most effective for the imposed initialization mismatch, whereas integration-
 432 based states generally require a later physical anchor.




Figure 8: Repeated noise robustness across six holdout cells. Bars show cell-macro ΔMAE and error bars show hierarchical 95% confidence intervals with random seeds nested within cells. HECM is most affected by high current noise, while most other channel-model combinations remain close to baseline relative to the between-cell uncertainty.

4.5. Missing samples and signal integrity

Sample loss is substantially more consequential than timing jitter. Periodic freezing of every 50th sample increases MAE by 0.0071, 0.0064, 0.0028, and 0.0011 for DM, HDM, HECM, and DD, respectively; random 2% loss produces corresponding increases of 0.0079, 0.0080, 0.0038, and 0.0026. By contrast, even the ± 0.9 s jitter case remains within approximately 6×10^{-4} SOC of baseline for every model, with confidence intervals generally spanning zero. DD therefore has the smallest sample-loss penalty, while DM and HDM retain their structural dependence on uninterrupted charge integration.

The one-hour burst dropout is a distinct recovery problem. Its macro ΔMAE is largest for HECM (0.0221), followed by HDM (0.0134), DD (0.0134), and DM (0.0067). After measurements resume, DD has the shortest mean recovery- or-censor time (4.13 h), then DM (5.13 h), HDM (5.67 h), and HECM (12.45 h). The hours-long response has model-specific origins: DM/HDM retain the unobserved charge deficit until a physical anchor, HECM must correct it through weak LFP voltage observability, and DD recovers as its rolling input context is replaced by valid samples. Figure 11 therefore reports the local transition in




Figure 9: Recovery after a 10% SOC-equivalent initialization mismatch. (a) Representative paired trajectories. (b) Difference between shifted and correctly initialized outputs with the common 2% threshold. (c) Stable recovery-or-censor time across the six holdout cells with hierarchical 95% confidence intervals.

450 addition to the global score.

451 4.6. Spike sensitivity

452 Single voltage spikes demonstrate why local and global metrics must be re-
 453 ported together. Full-window ΔMAE remains below 9×10^{-4} SOC for every
 454 estimator and is effectively zero for DM, HDM, and DD. Event-aligned analysis,
 455 however, identifies a reproducible DD transient: its mean spike-sample penalty
 456 across cells is approximately 7.7×10^{-3} SOC, compared with near-zero direct
 457 penalties for the other representatives. Figures 12 and 13 further show that the
 458 DD response varies across cells and SOH states and then decays after the event.
 459 This pattern is consistent with a one-sample voltage impulse entering both the
 460 absolute voltage and dU/dt channels and perturbing the rolling recurrent pre-
 461 diction. It is a mechanism-supported interpretation, not proof of a unique causal
 462 input channel because no channel-removal ablation was performed.




Figure 10: Six-cell signal-integrity and timing stress. Bars show cell-macro ΔMAE and hierarchical 95% confidence intervals for periodic and random sample loss and three timing-jitter levels. Sample loss clearly dominates timing jitter under the tested protocol.

4.7. Shared-SOH ablation

The primary comparison intentionally supplies one common causal LSTM-SOH trace to HDM, HECM, and DD. A paired ablation repeats selected scenarios with dataset reference SOH to quantify this shared-service confounder. Reference SOH reduces baseline MAE by 0.0363 for HDM and 0.0164 for HECM, but increases DD MAE by 0.0049 on average. The same direction persists for most noise, sample-loss, and dropout cases. At 3% current gain error, reference SOH still improves HDM and HECM but worsens DD by 0.0091. For temperature noise, the reference-minus-LSTM difference changes from its baseline value by only 0.0015 for HDM and less than 10^{-4} for HECM and DD, indicating that the incremental temperature-noise response is not dominated by the shared SOH service. This does not imply that a less accurate SOH estimate is intrinsically beneficial; rather, the models consume SOH differently and were calibrated with different feature interactions. The LSTM-SOH primary results therefore represent deployable pipeline performance, while reference-SOH runs are explanatory ablations rather than replacement rankings.




Figure 11: One-hour burst-dropout transition. (a) Baseline outputs (dashed) and dropout outputs (solid) around the shaded missing-input interval. (b) Absolute deviation from the paired no-dropout output after measurements resume. (c) Six-cell recovery-or-censor time under the common 2%/300-s criterion.

4.8. Robustness synthesis across representatives

Figure 14 condenses the complete measurement-only campaign and tests whether the nominal-accuracy ranking alone describes degraded-input behavior. It does not. Persistent gain error and missing information create the largest global penalties: current-gain sensitivity increases with magnitude for every representative, periodic and random sample loss particularly affect DM and HDM, and the one-hour dropout produces the largest penalty for HECM. Timing jitter remains nearly neutral, whereas voltage spikes illustrate the opposite limitation of global averaging: their full-window MAE is small even though event alignment exposes a distinct DD transient. DD has the lowest aggregate robustness penalty and is strongest under current-gain error and isolated sample loss in the tested campaign, but HECM re-anchors fastest after initialization mismatch and the integration-based representatives remain much cheaper to execute. The




Figure 12: Voltage-spike sensitivity across six cells. (a) Mean event-sample MAE penalty with hierarchical 95% confidence intervals. (b,c) Cell- and SOH-state-resolved penalties for HECM and DD, respectively; absent cell/state combinations are left blank. The event-aligned metric exposes the local DD response that is obscured in full-window Δ MAE.

central result is therefore a multi-dimensional ranking rather than one universal winner.

4.9. ADC quantization

The tested ADC steps ($\Delta I = 0.01$ A, $\Delta U = 0.005$ V, and $\Delta T = 0.5$ °C) preserve the local signal shape but affect the tested representatives differently. DM and HDM increase by 0.0037 and 0.0034 MAE, respectively. HECM changes by -0.0018 and DD by $+0.0002$; both confidence intervals overlap zero. Thus, quantization at these deliberately coarse steps is secondary to current gain error, sample loss, and burst dropout in this dataset, but it is not negligible for the integration-based models.

4.10. Target-hardware validation

The STM32 experiment first verifies numerical equivalence before comparing speed or memory. Across the six nominal cell sequences, mean hardware–software MAE remains below 4×10^{-4} percentage points for every SOC core and reaches 3.41×10^{-6} percentage points for DD (Figure 17). These differences are orders of magnitude below the estimator-to-dataset errors and show that the float32 C implementations reproduce their software references sufficiently closely for deployment measurements.




Figure 13: Event-aligned voltage-spike response. (a) Representative SOC trajectories around the injected spike; the dashed black line is reference SOC. (b) Model-wise 95th percentile peak excess error. (c) Mean excess absolute error aligned across repeated spikes, with the DD uncertainty envelope. The local response decays rapidly and is therefore weakly represented by full-window MAE.

510 The isolated DM and HDM SOC updates both have sub-microsecond median
 511 latency (0.171 and 0.165 μs), HECM requires 23.9 μs , and continuous-state DD
 512 requires 424.9 μs (Figure 18). At the nominal 1-Hz update rate, these medians
 513 occupy less than 0.043% of one sampling period, although this isolated deadline
 514 share is not a measurement of total BMS CPU load. The DD value in this
 515 cross-model comparison is the continuous-state deployment variant; exact 2024-
 516 sample rolling-window semantics are evaluated separately because recomputing
 517 the full sequence at every sample is much slower.

518 Release images occupy 27.0, 27.1, 470.1, and 116.1 KiB flash for DM, HDM,
 519 HECM, and DD, respectively. Static RAM is approximately 4.8 KiB for DM/HDM,
 520 4.9 KiB for HECM, and 69.8 KiB for rolling-window DD (Figure 19). HECM
 521 flash is dominated by its parameter surfaces, while the DD rolling window dom-




Figure 14: Cell-macro Δ MAE across the complete six-cell measurement-only campaign. Positive values denote degradation relative to each model’s paired baseline; small negative values reflect finite-window signed effects and are not interpreted as beneficial disturbances.




Figure 15: Decision-oriented synthesis for the four tested representatives. (a) Relative min-max-normalized accuracy, robustness, and recovery dimensions. (b) Composite scores under three illustrative 60/20/20 priority profiles. Scores summarize the present benchmark only and do not replace scenario-level evidence.

522 inates static RAM. These values describe the isolated firmware images and do
523 not include the shared SOH-LSTM service; static RAM also remains a lower
524 bound on peak runtime memory.

525 Figure 20 makes the DD deployment trade-off explicit. Exact rolling-window
526 inference preserves the trained software semantics but requires roughly 724 ms
527 median inference time, or 72.4% of the one-second period, and 69.8 KiB static
528 RAM. Continuous-state inference reduces latency to approximately 0.425 ms




Figure 16: ADC quantization. (a,b) Representative raw and quantized current and voltage. (c) Six-cell Δ MAE with hierarchical 95% confidence intervals and resulting quantized-signal MAE annotations.

529 and RAM to 6.5 KiB while retaining similar error on the six nominal sequences.
530 Periodically resetting that state has similar resource use but creates much larger
531 transient errors, with a maximum SOC error near 28 percentage points. Con-
532 tinuous state is therefore the practical hardware candidate, but it remains an
533 inference-semantics ablation rather than a drop-in replacement for the primary
534 rolling-window software benchmark.

535 5. Discussion

536 5.1. Structural interpretation of the representatives

537 Because one concrete implementation represents each broad estimator class,
538 the results support mechanism-level interpretation but not family-wide perfor-
539 mance claims. DM behaves as expected for a current-integration representative:
540 it has the largest nominal error and is sensitive to persistent gain error and miss-
541 ing samples because no voltage-feedback correction is available. Its practical
542 full-charge anchor limits unbounded drift, but the lifecycle experiment shows




Figure 17: Numerical equivalence of STM32 and software execution over six cell sequences. Bars show mean hardware–software MAE in percentage points, dots show cells, and whiskers show the cell range. The logarithmic scale emphasizes that all implementation differences are negligible relative to model error.

that the reset does not erase accumulated bias uniformly. HDM adapts the capacity denominator through learned SOH and can reduce capacity mismatch, yet it retains the same charge-integration core. Its similar response to current bias and missing charge information is therefore structural to this implementation rather than removed by SOH adaptation.

HECM combines current-driven state propagation with a voltage residual and recovers the imposed initial mismatch fastest. This correction capability does not make it universally robust: high current noise and burst dropout produce the largest HECM penalties because both state propagation and observer correction are affected when the causal input stream is corrupted or frozen. The primary campaign does not perturb the fixed HPPC-derived parameter surfaces, so the HECM current-gain result is conditional on this lookup-table quality rather than evidence of equal robustness under parameter misidentification. Its table dependence and flash footprint must therefore be balanced

On-Device Inference-Latency Distributions Across Six Cells




Figure 18: On-device SOC-core latency distributions across six cells and three replay rounds. Panels (a–d) show per-inference distributions; panel (e) compares minimum, 5th percentile, median, 95th percentile, and maximum on a logarithmic scale. DD denotes continuous-state inference in this cross-model latency comparison.

557 against its rapid initialization correction.

558 DD has the lowest nominal cell-macro error and the best adverse current-
559 gain score in this comparison. It is also comparatively resilient to isolated
560 sample loss and recovers fastest after burst dropout. Conversely, event-aligned
561 voltage spikes expose a recurrent transient that full-window MAE hides. This




Figure 19: Measured release-image flash and static-RAM footprints for the isolated STM32 SOC implementations. Flash and RAM are firmware properties and are therefore reported once per implementation rather than as artificial cell-dependent distributions.




Figure 20: DD inference-mode trade-off on the STM32. The panels compare maximum dataset-SOC error, median latency, and firmware memory for exact rolling-window, continuous-state, and periodically reset recurrent execution. Cell points show that periodic resets, not continuous forwarding, create the dominant accuracy loss.

562 vulnerability is consistent with the selected voltage and derivative features and
 563 rolling recurrent context. It is specific to the frozen GRU, feature set, training
 564 data, and inference semantics, not a universal limitation of recurrent or data-
 565 driven SOC estimation.

566 5.2. Why nominal accuracy is insufficient

567 The central finding is not that nominal accuracy and robustness are un-
 568 related: the tested DD representative leads both the clean-data ranking and

the aggregate robustness score. Rather, nominal MAE alone fails to reveal the conditions under which that ranking changes or becomes incomplete. It does not expose the DD voltage-spike transient, HECM’s substantially faster correction of initialization mismatch, the different recovery order after dropout, the shared-SOH feature coupling, or the orders-of-magnitude difference in execution cost. No single scalar captures these deployment trade-offs. Selection should therefore begin with the relevant disturbance and recovery requirements and then consider latency, memory, and state-service architecture, rather than treating baseline accuracy as a sufficient proxy.

5.3. Implications for BMS deployment

Among the tested representatives, DD is the strongest aggregate accuracy/robustness candidate if its recurrent implementation cost and voltage-spike response are acceptable. HECM is preferable when rapid correction after state mismatch dominates and a larger flash image is available. DM and HDM remain attractive when deterministic sub-microsecond execution and small memory dominate, but persistent current calibration, reliable sampling, and full-charge anchoring then become stronger system-level requirements. These are implementation-level recommendations, not claims that one estimator family is intrinsically superior.

The STM32 results narrow the gap between algorithmic and deployment claims. All four isolated SOC cores meet the one-second deadline with the tested implementation, but their margins differ by more than six orders of magnitude between DM/HDM and exact rolling-window DD. Continuous-state DD removes most of that cost, at the price of changing inference semantics. The hardware test does not include the shared SOH-LSTM, peak stack, energy, pack communication, or concurrent BMS tasks; therefore it supports SOC-core feasibility, not complete BMS schedulability. Table 3 summarizes the evidence without hiding these boundaries.

Figure 15 complements the recommendation map with relative scores, not additional measurements. Lower-is-better inputs are min–max normalized within the four tested implementations. Accuracy combines baseline MAE, RMSE,

Table 3: Decision-oriented recommendation map derived from the benchmark results. The table identifies the strongest tested representative when one deployment priority dominates the design decision.

Primary deployment priority	Selected representative	Main supporting benchmark evidence	Main trade-off to consider
Lowest nominal SOC error	DD	Lowest six-cell macro MAE (0.0258) and RMSE (0.0300)	Rolling-window execution is the most computationally expensive implementation
Fastest recovery from initialization mismatch	HECM	Mean stable recovery of 1.03 h under the common criterion	Highest flash footprint and slowest post-dropout recovery
Lowest adverse current-gain penalty	DD	Δ MAE 0.0038 at 3%, below HECM, HDM, and DM	Local voltage-spike transient remains visible
Lowest periodic/random sample-loss penalty	DD	Δ MAE 0.0011/0.0026	One-hour burst dropout still increases MAE by 0.0134
Lowest deterministic compute cost	DM/HDM	Sub-microsecond median SOC-core latency and approximately 27 KiB flash	Integration structure is sensitive to gain error and missing charge information
Practical fast neural deployment	DD continuous state	Approximately 0.425 ms median and 6.5 KiB static RAM	Different recurrent semantics from the primary rolling-window benchmark

and 95th-percentile error; robustness combines cross-scenario penalties; and recovery combines initialization and burst-dropout behavior. The three profiles assign 60% to their named dimension and 20% to each remaining dimension. DD scores highest in all three illustrative profiles (0.933, 0.842, and 0.890), followed by HECM, HDM, and DM. This consistency reflects the selected metrics and normalization; it neither removes scenario-specific weaknesses nor supports extrapolation from these representatives to their complete model families.

5.4. Limitations and future extensions

Only one frozen implementation represents each broad estimator family; different reset logic, ECM order and parameterization, neural architecture, input features, training data, or inference semantics can change the ranking. The study is further limited to six holdout cells from one controlled LFP dataset and therefore cannot establish transfer to NMC/NCA chemistry, other form factors, pack imbalance, other thermal environments, or field duty cycles. The load groups are descriptive strata with $n = 2$ low-load, $n = 3$ middle-load, and one exploratory high-load cell; they are not randomized treatments. C27 contains no mid-life or aged window, and missing states are not imputed. With six

independent cells, confidence intervals and effect magnitudes are more informative than dichotomized significance decisions.

Only measurement, timing, availability, and initialization channels are manipulated. Cell/pack parameter mismatch, ECM aging drift, hysteresis error, thermal gradients, sensor cross-sensitivity, balancing and contactor faults, communication-protocol failures, numerical faults, and deadline interference remain outside the campaign. DD initialization is implemented through an SOC-equivalent Q_c offset rather than direct hidden-state corruption. Abrupt capacity-loss events are absent, so the hourly SOH service is evaluated only for gradual aging. The voltage-spike mechanism is supported by event alignment but not proven by an input-channel ablation. HECM results remain conditional on the imported HPPC lookup tables, which were fixed rather than reidentified for each holdout cell.

The hardware experiment isolates SOC cores and passes SOH as an external causal trace. It does not measure end-to-end SOH latency, peak stack, energy, CPU utilization under concurrent BMS tasks, or pack-level closed-loop behavior. Future work should therefore prioritize peak-memory and energy instrumentation, scheduling tests with the SOH service, clustered-fault experiments, hidden-state ablations, and independent validation on additional chemistries and duty cycles.

6. Conclusion

This study asked whether nominal SOC accuracy is sufficient to select an estimator for disturbed sensing and embedded deployment. For the four tested representatives, it is not. DD provides the lowest nominal error and the strongest aggregate robustness, but its local voltage-spike response and rolling-window cost are invisible in a clean-data MAE ranking. HECM restores an incorrect initial state fastest but is less tolerant of current noise and prolonged dropout and requires the largest flash image. DM and HDM execute with minimal latency and memory, yet remain most dependent on accurate, continuous current

integration and physical re-anchoring.

The practical outcome is therefore a selection framework rather than a universal class ranking: accuracy characterizes nominal prediction, robustness identifies disturbance-specific degradation, recovery describes re-synchronization, and hardware measurements determine whether the trained execution semantics remain practical. The STM32 tests confirm feasibility of all isolated SOC cores while exposing substantial implementation trade-offs. These conclusions apply to the selected Coulomb-counting, SOH-adaptive, ECM-EKF, and GRU representatives on the tested LFP aging data; broader claims require additional implementations, chemistries, operating profiles, and complete SOC-SOH hardware scheduling.

Appendix A. Holdout Coverage and Frozen Protocol




Figure A.21: Quantitative coverage of the six independent holdout cells. Panels show observed SOH, the descriptive load-class assignment, thermal envelope, and full-trajectory SOH-state coverage. High-load evidence consists only of C29 and is exploratory.




Figure A.22: Complete 19-case JES2 matrix. Red markers identify the manipulated measurement, timing, availability, or initialization channel; blue-gray markers denote repeated stochastic cases; rose markers identify paired reference-SOH ablations.




Figure A.23: Frozen representative-window protocol. Sixteen cell/SOH windows are selected using measured-feature medoids only. Each scored 24-h window follows a 192-h causal SOH context; the dropout case extends evaluation to 48 h.




Figure A.24: DD latency distributions over the six hardware cell sequences. Exact rolling-window execution is shown separately from continuous-state and periodically reset execution because their scales and recurrent semantics differ substantially.

Table B.4: Six-cell baseline cell-macro statistics with hierarchical 95% confidence intervals.

Estimator	MAE [SOC]	RMSE [SOC]	P95 error [SOC]
DM	0.0678 [0.0491, 0.0809]	0.0733 [0.0524, 0.0873]	0.1086 [0.0754, 0.1307]
HDM	0.0405 [0.0297, 0.0506]	0.0438 [0.0317, 0.0555]	0.0648 [0.0439, 0.0845]
HECM	0.0335 [0.0246, 0.0429]	0.0389 [0.0289, 0.0496]	0.0646 [0.0467, 0.0812]
DD	0.0258 [0.0189, 0.0323]	0.0300 [0.0224, 0.0367]	0.0508 [0.0375, 0.0618]

Table B.5: Selected cell-macro scenario penalties relative to paired baseline. Current-gain values use the adverse sign from each matched $\pm 3\%$ pair.

Scenario	DM	HDM	HECM	DD
Current noise, $\sigma_I = 0.10$ A	-0.0009	+0.0000	+0.0056	-0.0000
Current gain, adverse 3%	+0.0107	+0.0100	+0.0059	+0.0038
ADC quantization	+0.0037	+0.0034	-0.0018	+0.0002
Random missing samples, 2%	+0.0079	+0.0080	+0.0038	+0.0026
Timing jitter, ± 0.9 s	-0.0003	+0.0003	+0.0006	+0.0001
Burst dropout, 1 h	+0.0067	+0.0134	+0.0221	+0.0134
Voltage spikes, ± 0.20 V	+0.0000	-0.0003	+0.0009	-0.0001

Table B.6: Recovery and target-hardware summary. Recovery values are equal-weight cell means.

Metric	DM	HDM	HECM	DD
Initial stable recovery [h]	22.28	22.28	1.03	5.57
Burst recovery/censor time [h]	5.13	5.67	12.45	4.13
Median SOC-core latency [μ s]	0.171	0.165	23.9	424.9 ^a
Flash [KiB]	27.0	27.1	470.1	116.1
Static RAM [KiB]	4.8	4.8	4.9	69.8

^aContinuous-state DD; exact rolling-window median is approximately 724 ms.

658 **Funding**

659 Funding details are omitted for double-blind review.

660 **Acknowledgements**

661 Acknowledgements are omitted for double-blind review.

662 **CRediT authorship contribution statement**

663 The author contribution statement is provided in the non-blinded submission
664 materials.

665 **Declaration of competing interest**

666 The authors declare that they have no known competing financial interests or
667 personal relationships that could have appeared to influence the work reported
668 in this paper.

669 **Declaration of generative AI and AI-assisted technologies**
670 **in the manuscript preparation process**

671 During the preparation of this work the author(s) used ChatGPT, DeepL,
672 and Grammarly in order to improve the linguistic quality and style of the
673 manuscript. After using these tools, the author(s) reviewed and edited the con-
674 tent as needed and take(s) full responsibility for the content of the published
675 article.

676 **Data availability**

677 The dataset is publicly available through its cited TU Berlin repository
678 record [28]; code and results are maintained in a public GitHub repository. A
679 versioned release will archive the fixed manifests and seeds, frozen model/scaler/ECM
680 artifacts and hashes, causal runners, run summaries, analysis and plotting
681 scripts, and environment needed to rebuild all reported results.

References

- [1] P. Shrivastava, P. A. Naidu, S. Sharma, B. K. Panigrahi, A. Garg, Review on technological advancement of lithium-ion battery states estimation methods for electric vehicle applications, *Journal of Energy Storage* 64 (2023) 107159. doi:10.1016/j.est.2023.107159.
- [2] F. Guo, L. D. Couto, G. Mulder, K. Trad, G. Hu, O. Capron, K. Haghverdi, A systematic review of electrochemical model-based lithium-ion battery state estimation in battery management systems, *Journal of Energy Storage* 101 (2024) 113850. doi:10.1016/j.est.2024.113850.
- [3] J. Yao, J. Kowal, Towards a smarter battery management system: A critical review on deep learning-based state of charge estimation of lithium-ion batteries, *Energy and AI* 21 (2025) 100585. doi:10.1016/j.egyai.2025.100585.
- [4] X. Li, J. Jiang, C. Zhang, L. Y. Wang, L. Zheng, Robustness of SOC Estimation Algorithms for EV Lithium-Ion Batteries against Modeling Errors and Measurement Noise, *Mathematical Problems in Engineering* 2015 (2015) 1–14. doi:10.1155/2015/719490.
- [5] F. Naseri, C. Barbu, A. Jensen, P. Setoodeh, J. R. Baena, N. Hevele, E. Schaltz, Toward Reliable Wireless BMS: Robust Battery State Estimation Under Noise and Uncertainty, *IEEE Transactions on Vehicular Technology* (2025) 1–12doi:10.1109/TVT.2025.3638613.
- [6] F. Berger, D. Joest, E. Barbers, K. Quade, Z. Wu, D. U. Sauer, P. Dechent, Benchmarking battery management system algorithms - Requirements, scenarios and validation for automotive applications, *eTransportation* 22 (2024) 100355. doi:10.1016/j.etrans.2024.100355.
- [7] Z. A. Khan, P. Shrivastava, S. M. Amrr, S. Mekhilef, A. A. Algethami, M. Seyedmahmoudian, A. Stojcevski, A Comparative Study on Different

- 709 Online State of Charge Estimation Algorithms for Lithium-Ion Batteries,
710 Sustainability 14 (12) (2022) 7412. doi:10.3390/su14127412.
- 711 [8] Y. Gao, G. L. Plett, G. Fan, X. Zhang, Enhanced state-of-charge estimation
712 of LiFePO₄ batteries using an augmented physics-based model, Journal
713 of Power Sources 544 (2022) 231889. doi:10.1016/j.jpowsour.2022.
714 231889.
- 715 [9] Y. Gao, K. Liu, C. Zhu, X. Zhang, D. Zhang, Co-Estimation of State-of-
716 Charge and State-of- Health for Lithium-Ion Batteries Using an Enhanced
717 Electrochemical Model, IEEE Transactions on Industrial Electronics 69 (3)
718 (2022) 2684–2696. doi:10.1109/TIE.2021.3066946.
- 719 [10] J. Shen, W. Ma, J. Xiong, X. Shu, Y. Zhang, Z. Chen, Y. Liu, Alterna-
720 tive combined co-estimation of state of charge and capacity for lithium-
721 ion batteries in wide temperature scope, Energy 244 (2022) 123236. doi:
722 10.1016/j.energy.2022.123236.
- 723 [11] C. Wang, N. Cui, Z. Cui, H. Yuan, C. Zhang, Fusion estimation of lithium-
724 ion battery state of charge and state of health considering the effect of
725 temperature, Journal of Energy Storage 53 (2022) 105075. doi:10.1016/
726 j.est.2022.105075.
- 727 [12] P. Mondal, D. Bhavsar, K. Mittal, M. Mittal, Estimating State-of-Charge
728 in Lithium-Ion Batteries Through Deep Learning Techniques: A Com-
729 parative Evaluation, IEEE Access 12 (2024) 78773–78786. doi:10.1109/
730 ACCESS.2024.3408220.
- 731 [13] D. P. Pau, A. Aniballi, Tiny Machine Learning Battery State-of-Charge
732 Estimation Hardware Accelerated, Applied Sciences 14 (14) (2024) 6240.
733 doi:10.3390/app14146240.
- 734 [14] S. Giazitzis, M. Sakwa, S. Leva, E. Ogliari, S. Badha, F. Rosetti, A
735 Case Study of a Tiny Machine Learning Application for Battery State-

- of-Charge Estimation, Electronics 13 (10) (2024) 1964. doi:10.3390/electronics13101964.
- [15] Y. Mazzi, H. Ben Sassi, A. Gaga, F. Errahimi, State of charge estimation of an electric vehicle's battery using tiny neural network embedded on small microcontroller units, International Journal of Energy Research 46 (6) (2022) 8102–8119. doi:10.1002/er.7713.
- [16] M. A. Kamali, A. C. Caliwag, W. Lim, Novel SOH Estimation of Lithium-Ion Batteries for Real-Time Embedded Applications, IEEE Embedded Systems Letters 13 (4) (2021) 206–209. doi:10.1109/LES.2021.3078443.
- [17] B. Gou, Y. Xu, X. Feng, An Ensemble Learning-Based Data-Driven Method for Online State-of-Health Estimation of Lithium-Ion Batteries, IEEE Transactions on Transportation Electrification 7 (2) (2021) 422–436. doi:10.1109/TTE.2020.3029295.
- [18] W. Li, N. Sengupta, P. Dechent, D. Howey, A. Annaswamy, D. U. Sauer, Online capacity estimation of lithium-ion batteries with deep long short-term memory networks, Journal of Power Sources 482 (2021) 228863. doi:10.1016/j.jpowsour.2020.228863.
- [19] F. Rzepka, P. Hematty, M. Schmitz, J. Kowal, Neural Network Architecture for Determining the Aging of Stationary Storage Systems in Smart Grids, Energies 16 (17) (2023) 6103. doi:10.3390/en16176103.
- [20] BQ79616 16-Series Battery Monitor, Balancer, and Integrated Hardware Protector (2023).
- [21] Design Considerations of Current Sensing With BQ769x2 Family for Battery Management System (2022).
- [22] TMCS1108 3% Functional Isolation Hall-Effect Current Sensor With ± 100 -V Working Voltage (2024).

- 762 [23] TLE4972 high precision coreless current sensor for automotive applications
763 (2023).
- 764 [24] O. Aiello, Electromagnetic Susceptibility of Battery Management Systems’
765 ICs for Electric Vehicles: Experimental Study, *Electronics* 9 (3) (2020) 510.
766 doi:10.3390/electronics9030510.
- 767 [25] M. Liu, D. Geißler, S. Bian, B. Zhou, P. Lukowicz, Assessing the Impact of
768 Sampling Irregularity in Time Series Data: Human Activity Recognition
769 As A Case Study (Jan. 2025). arXiv:2501.15330, doi:10.48550/arXiv.
770 2501.15330.
- 771 [26] 10s battery pack monitoring, balancing, and comprehensive protection, 50-
772 A discharge reference design (2024).
- 773 [27] A. Aqachtoul, K. Karam, A. Elamrani, Y. Alidrissi, M. Najib, A. Rochd,
774 M. Bakhouya, MGFARM: An IoT and Edge-Driven Framework for Smart
775 and Sustainable Agricultural Practices, in: Y. Mejdoub, A. Elamri, M. Kar-
776 douchi (Eds.), *Connected Objects, Artificial Intelligence, Telecommunica-*
777 *tions and Electronics Engineering*, Vol. 1584, Springer Nature Switzerland,
778 Cham, 2026, pp. 419–425. doi:10.1007/978-3-032-01536-5_64.
- 779 [28] F. Rzepka, Design of Experiments for Cyclic Aging of LFP 18650 Cells: Im-
780 pact of Charge/Discharge C-Rates, Depth of Discharge, and Temperature
781 on Battery Degradation (Nov. 2025). doi:10.14279/DEPOSITONCE-24800.
- 782 [29] K. Siebertz, D. Van Bebber, T. Hochkirchen, *Statistische Versuchspla-*
783 *nung*, Springer Berlin Heidelberg, Berlin, Heidelberg, 2017. doi:10.1007/
784 978-3-662-55743-3.